

**HW2 deadline postponed to next Thu, Oct 31!**

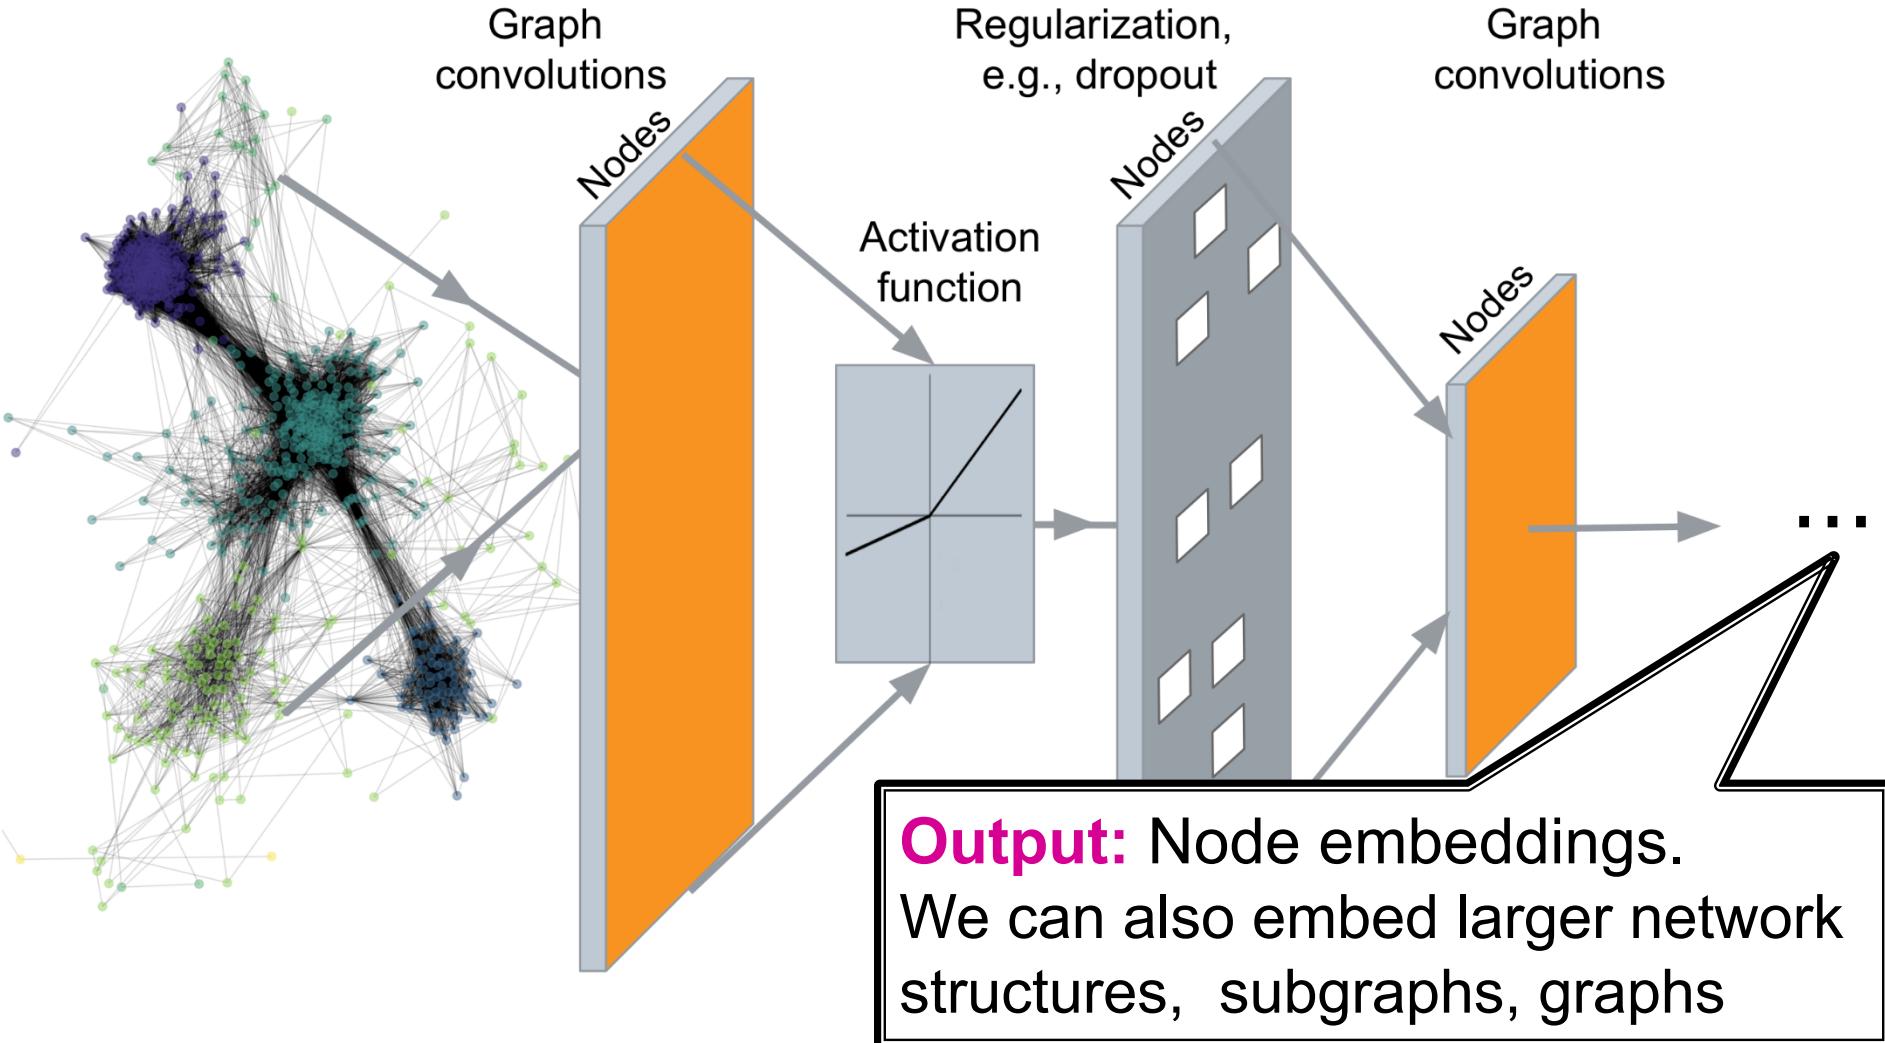
We are releasing an **improved version** of the  
starter code for HW2.Q4 -- keep an eye on Piazza!

# Deep Generative Models for Graphs

CS224W: Machine Learning with Graphs  
Jure Leskovec, Jiaxuan You, Stanford University  
<http://cs224w.stanford.edu>

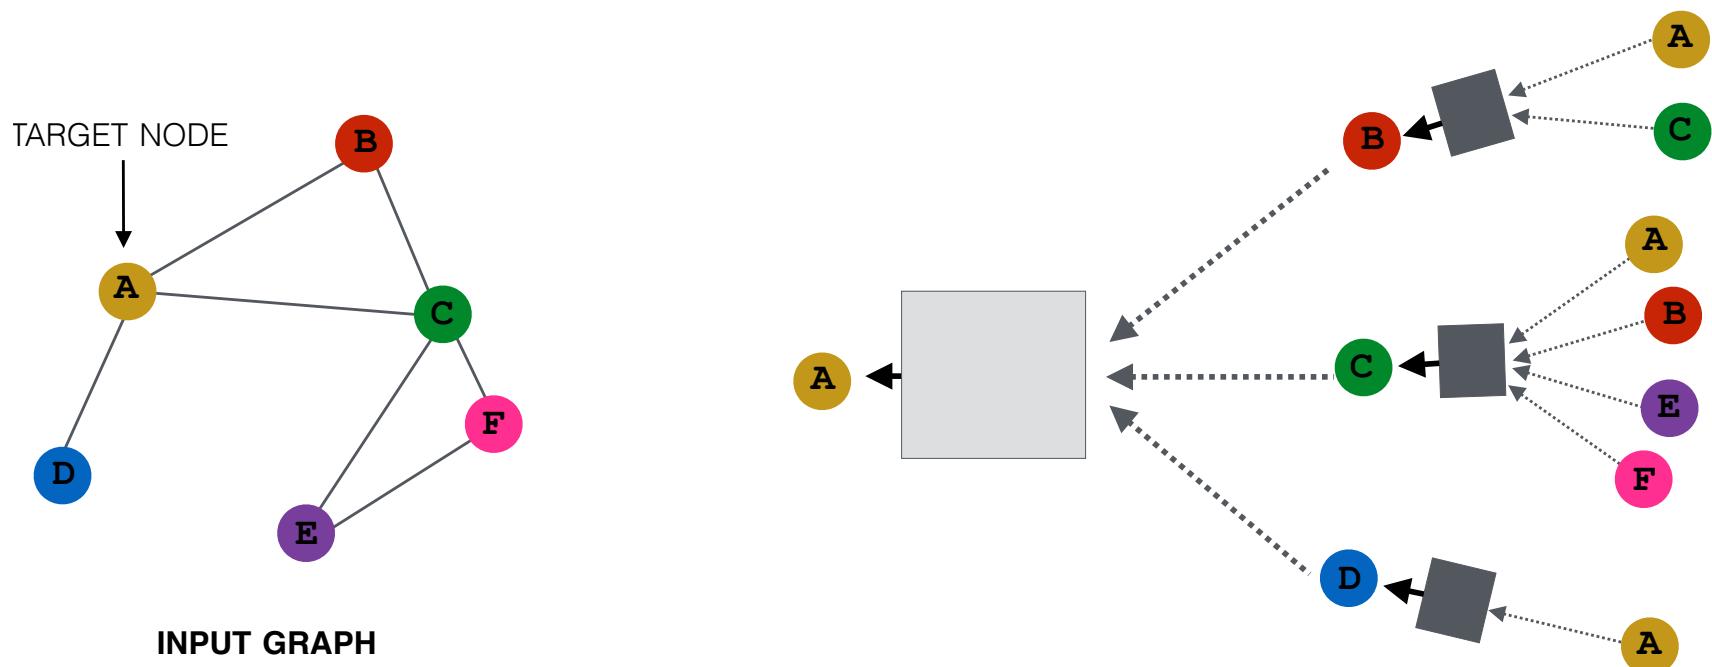


# Deep Graph Encoders



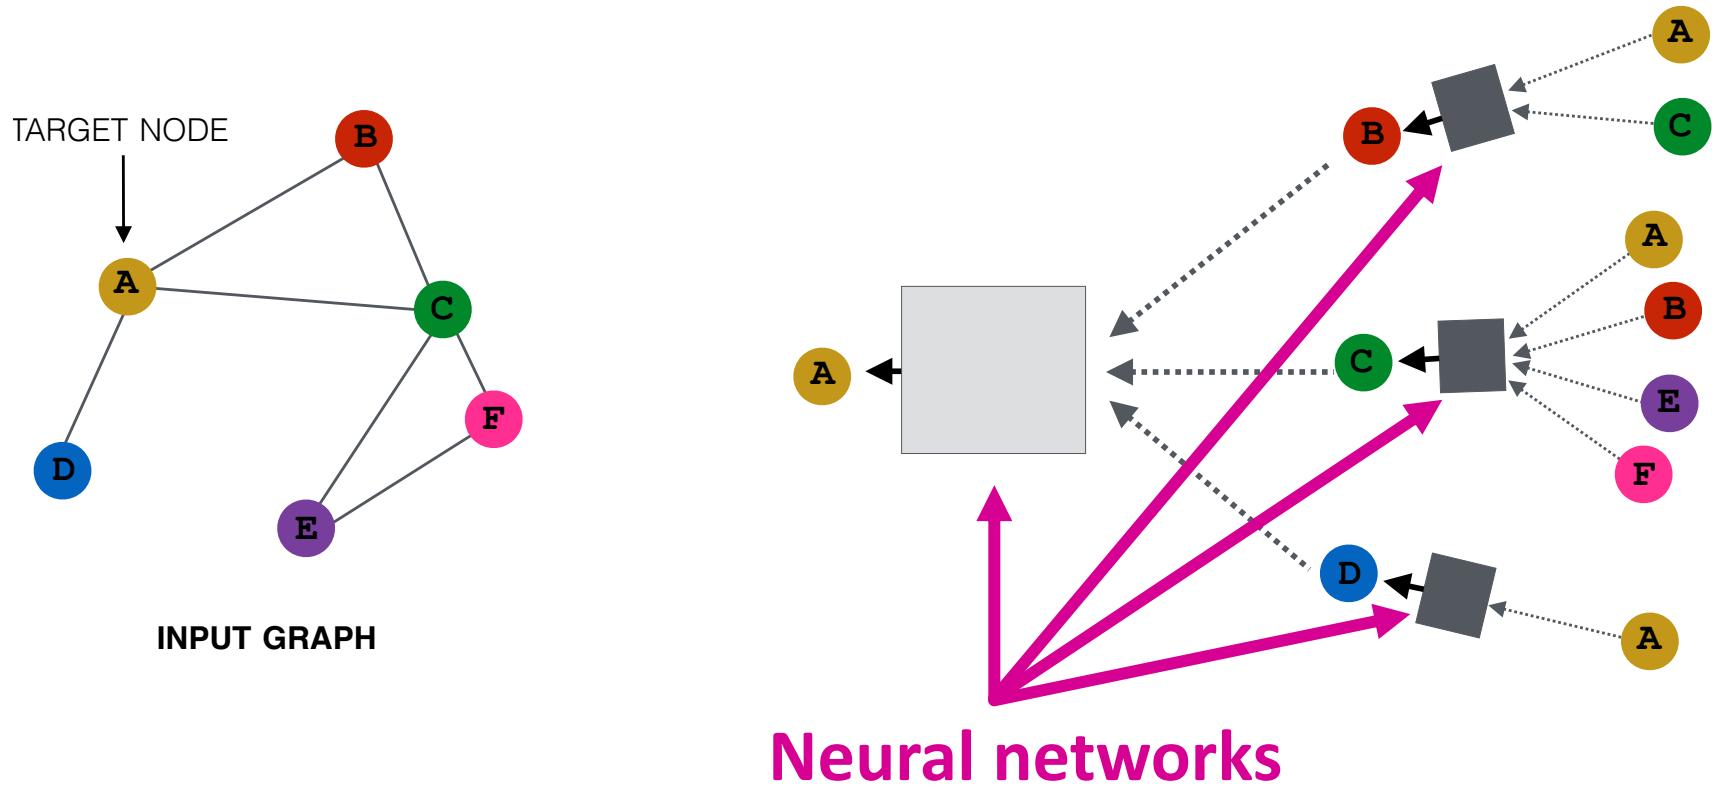
# Idea: Aggregate Neighbors

- **Key idea:** Generate node embeddings based on **local network neighborhoods**



# Idea: Aggregate Neighbors

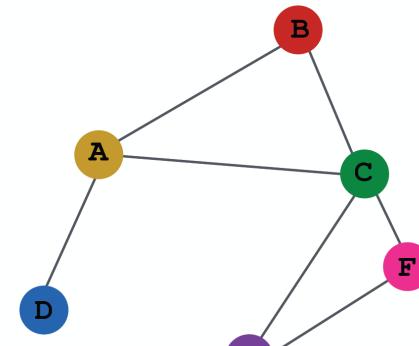
- **Intuition:** Nodes aggregate information from their neighbors using neural networks



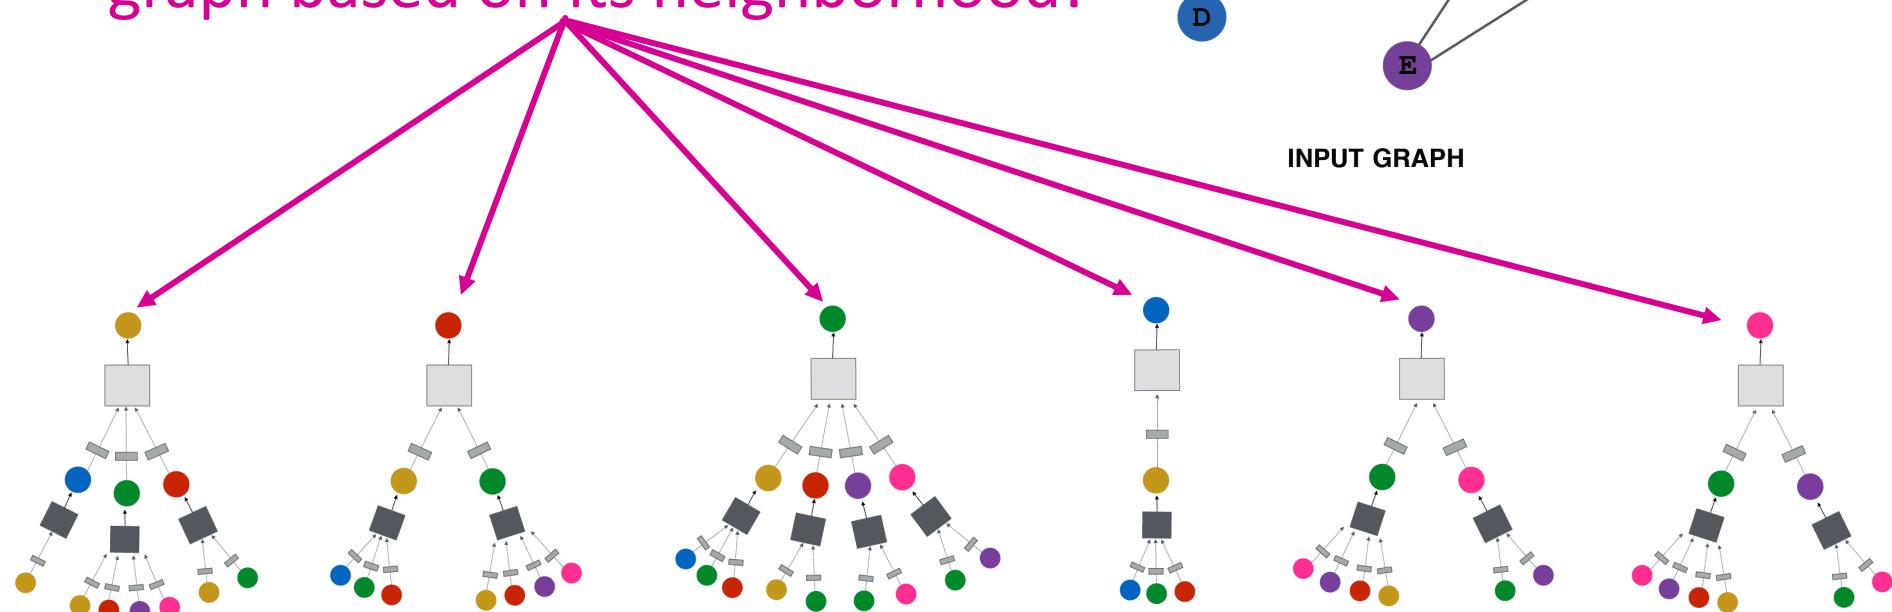
# Idea: Aggregate Neighbors

- **Intuition:** Network neighborhood defines a computation graph

Every node defines a computation graph based on its neighborhood!



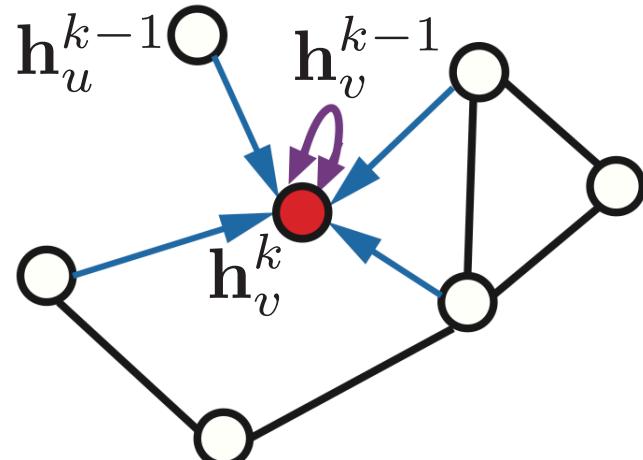
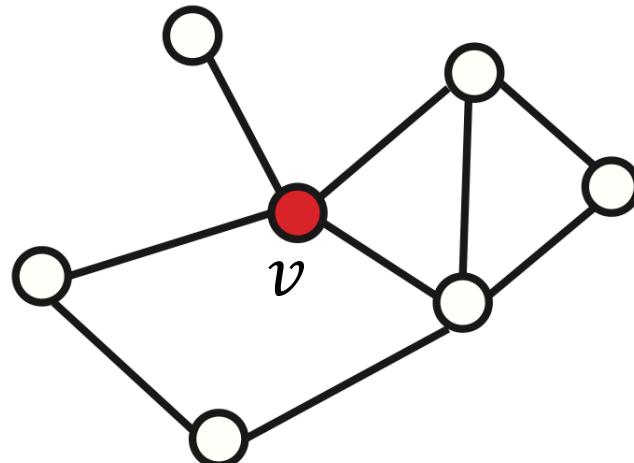
INPUT GRAPH



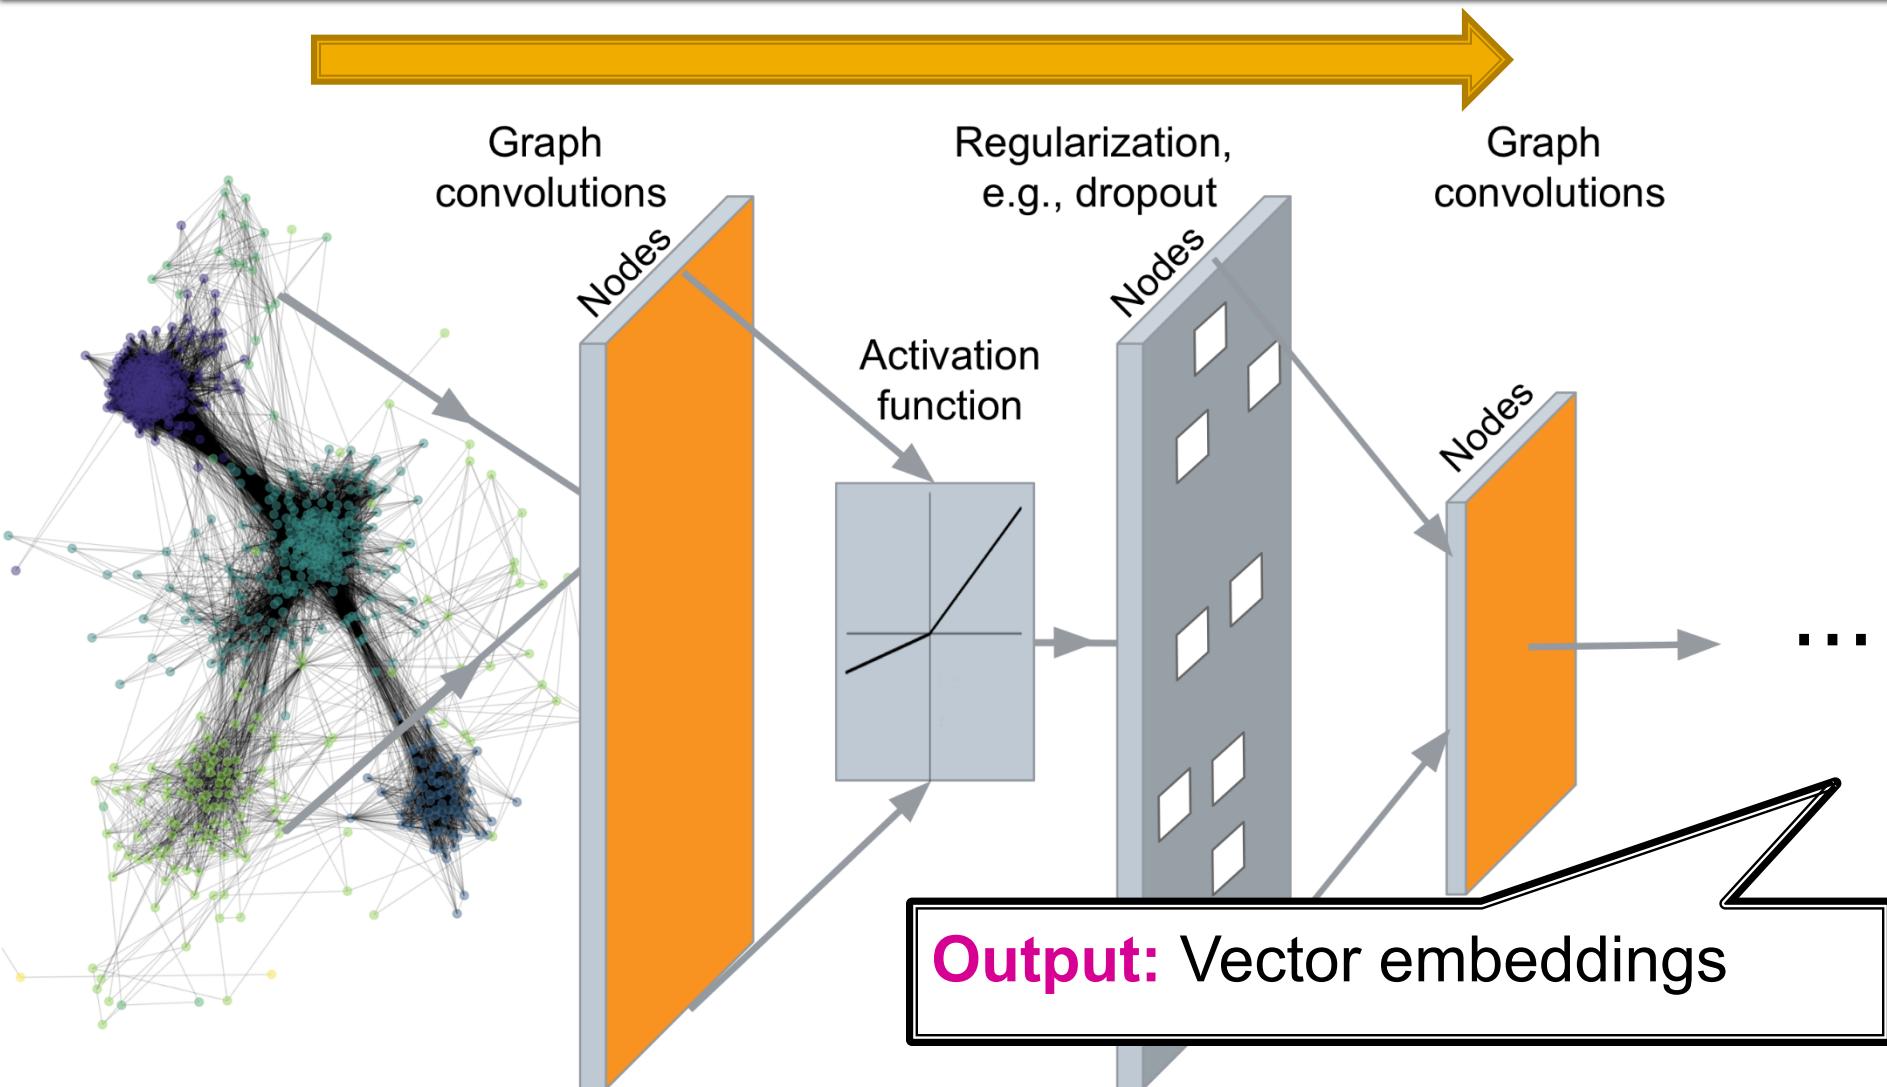
# Recap: GCN, GraphSAGE

**Key idea:** Generate node embeddings based on local network neighborhoods

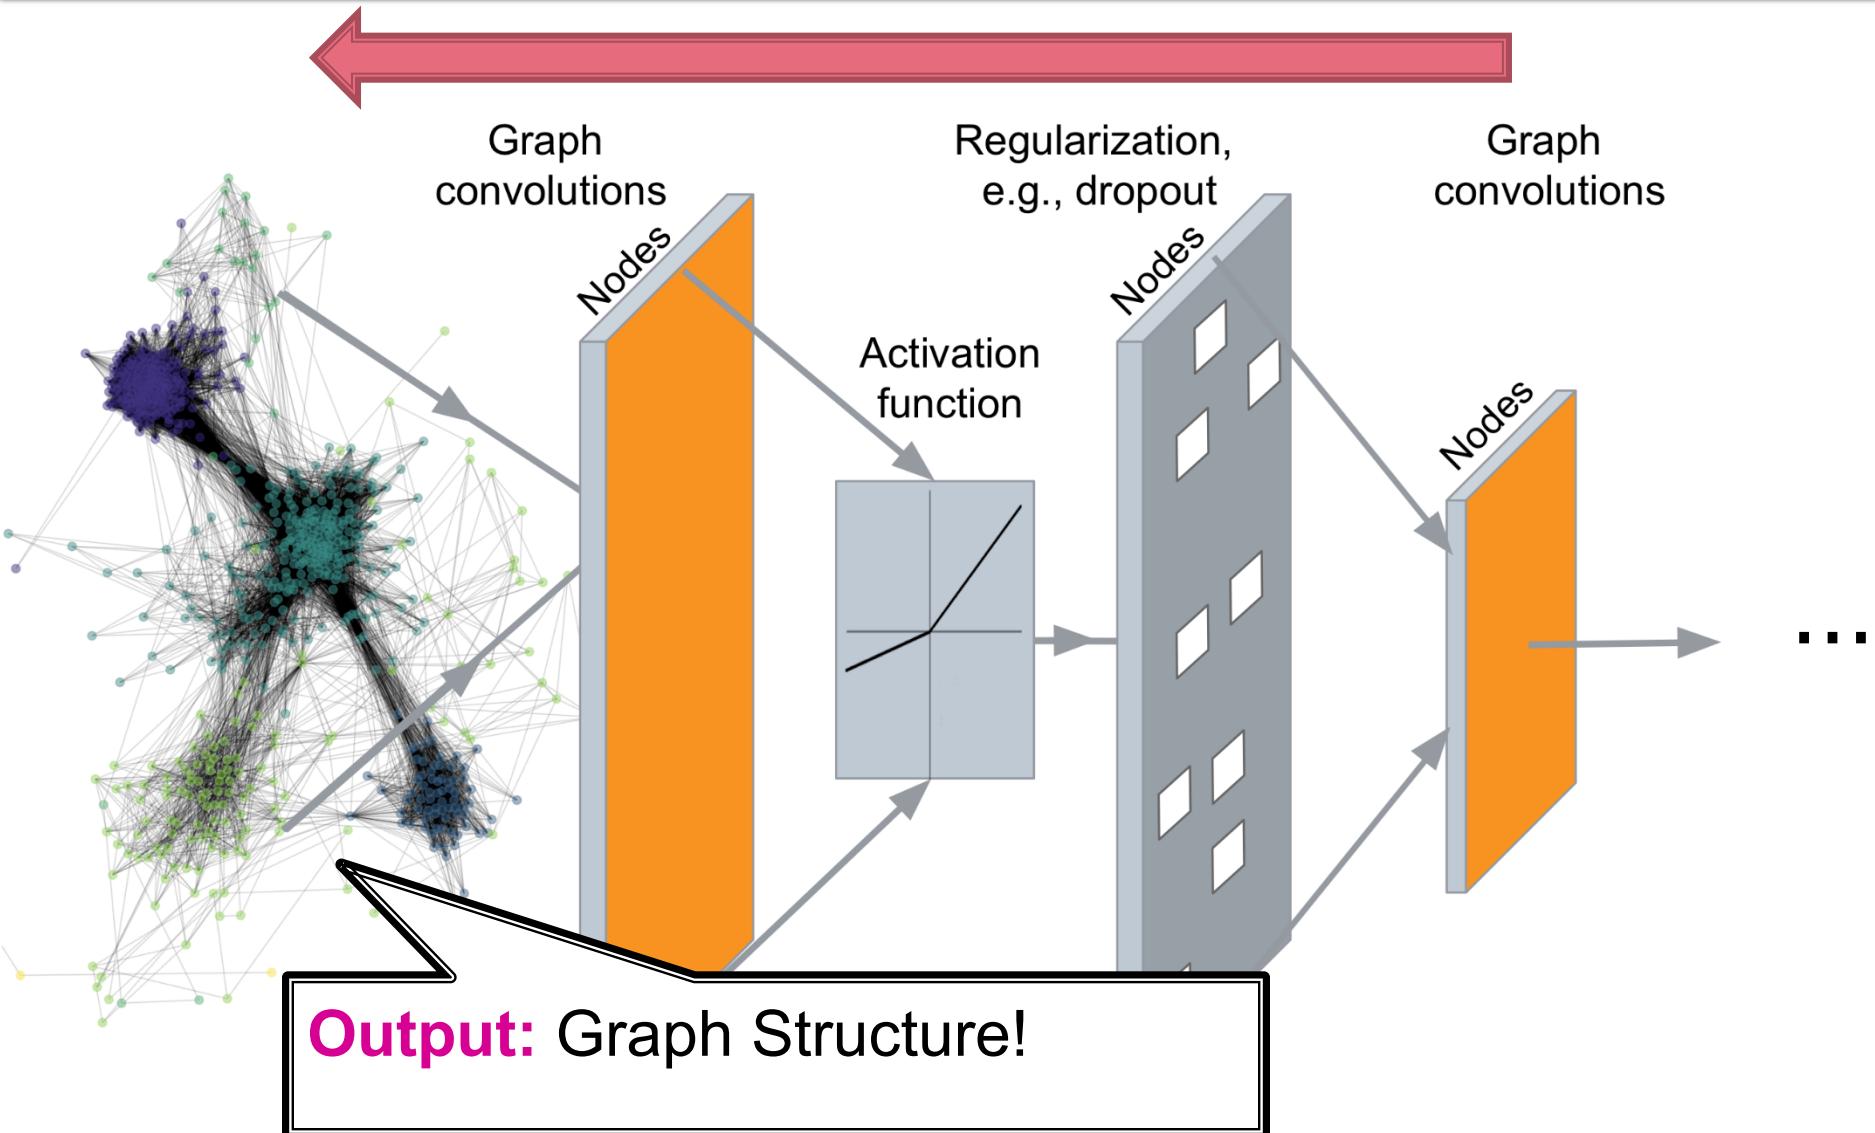
- Nodes aggregate “messages” from their neighbors using neural networks
- **Graph Convolutional Neural Networks:**
  - **Basic variant:** Average neighborhood information and stack neural networks
- **GraphSAGE:**
  - Generalized neighborhood aggregation



# Deep Graph Encoders



# Today: Deep Graph Decoders



# Outline of Today's Lecture

**1. Problem of Graph Generation**



**2. ML Basics for Graph Generation**

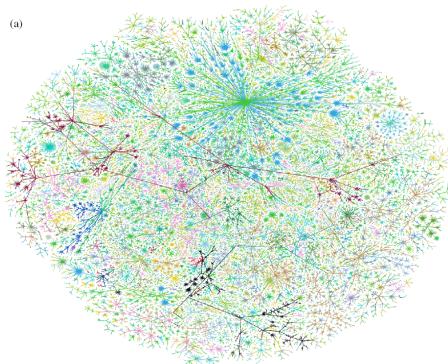
**3. GraphRNN**

**4. Applications and Open Questions**

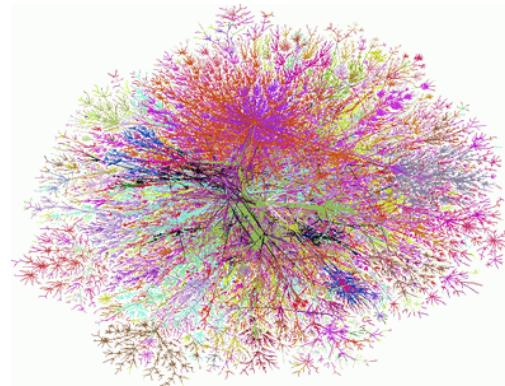
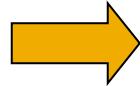
# The Problem: Graph Generation

# The Problem: Graph Generation

- We want to generate realistic graphs



Given a large  
real graph



Generate a  
synthetic graph

- What is a good model?
- How can we fit the model and generate the graph using it?

# Why is it Important?

- *Generation* – Gives insight into the graph formation process
- *Anomaly detection* – abnormal behavior, evolution
- *Predictions* – predicting future from the past
- *Simulations* of novel graph structures
- *Graph completion* – many graphs are partially observed
- *"What if"* scenarios

# Graph Generation Tasks

## Task 1: Realistic graph generation

- Generate graphs that are similar to a given set of graphs **[Focus of this lecture]**

## Task 2: Goal-directed graph generation

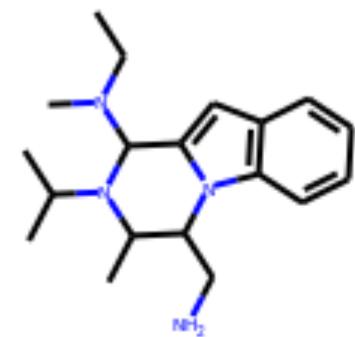
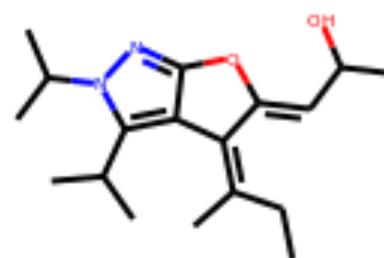
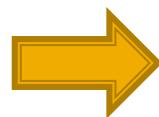
- Generate graphs that optimize given objectives/constraints
  - Drug molecule generation/optimization

# Why is it Interesting?

## Drug discovery

- Discover highly drug-like molecules

Graph  
generative  
model

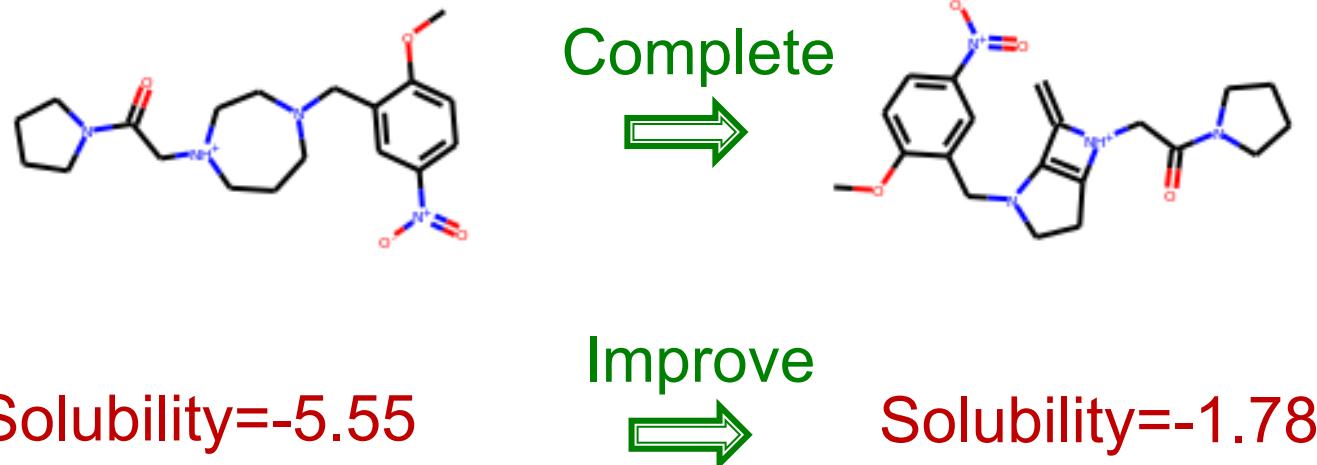


$\text{drug\_likeness} = 0.94$

# Why is it Interesting?

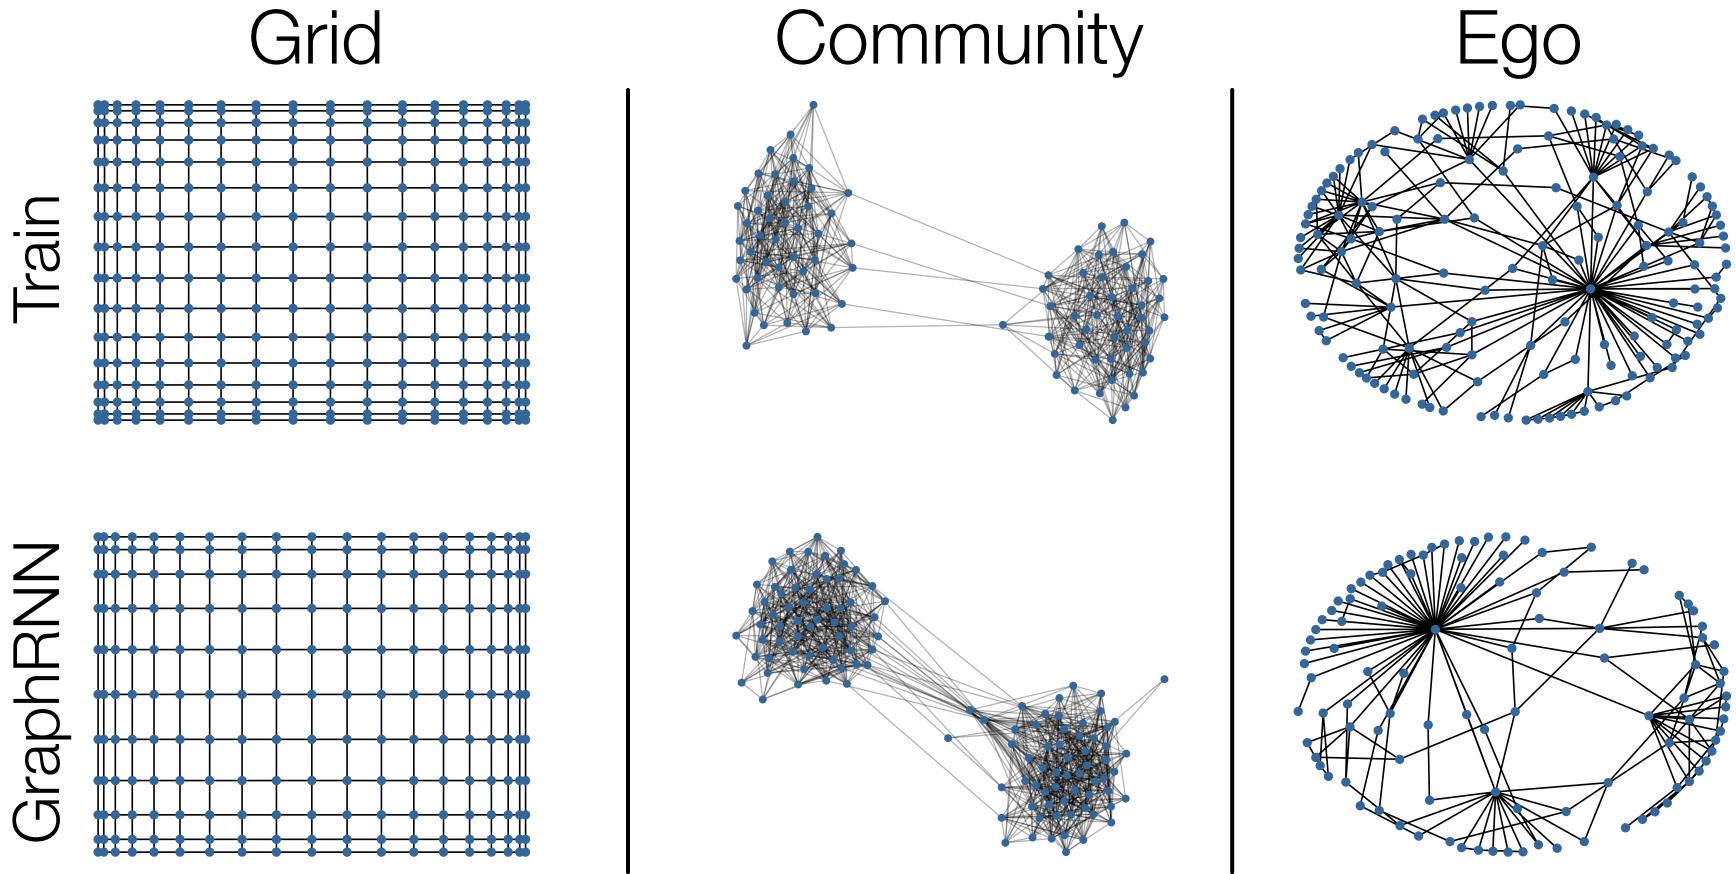
## Drug discovery

- Complete an existing molecule to optimize a desired property



# Why is it Interesting?

## Discovering novel structures

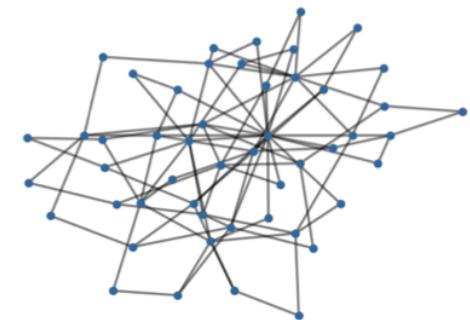


# Why is it Interesting?

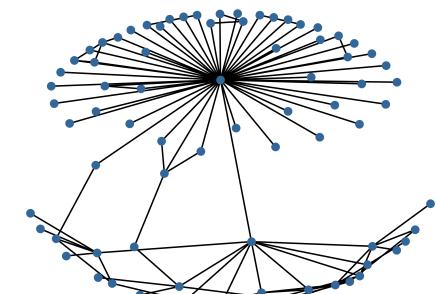
## Network Science

- Null models for realistic networks

Barabasi\_Albert( $n=50$ ,  $m=2$ ) ~



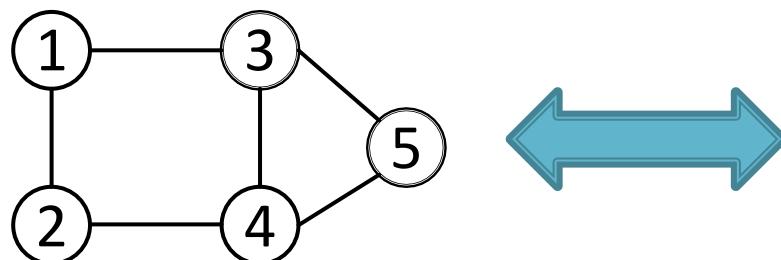
NeuralNet\_X( $n=50$ ,  $p=3$ ,  $q=5$ ) ~



# Why is it Hard?

## ■ Large and variable output space

- For  $n$  nodes we need to generate  $n^2$  values
- Graph size (nodes, edges) varies



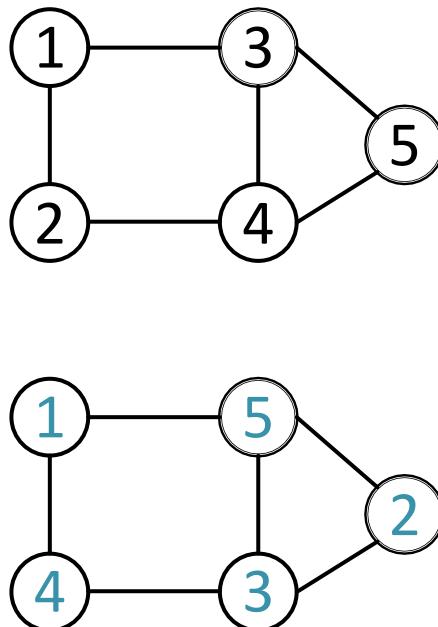
0	1	1	0	0
1	0	0	1	0
1	0	0	1	1
0	1	1	0	1
0	0	1	1	0

5 nodes: 25 values

1K nodes: 1M values

# Why is it Hard?

- **Non-unique representations:**
  - $n$ -node graph can be represented in  $n!$  ways
  - Hard to compute/optimize objective functions (e.g., reconstruction error)



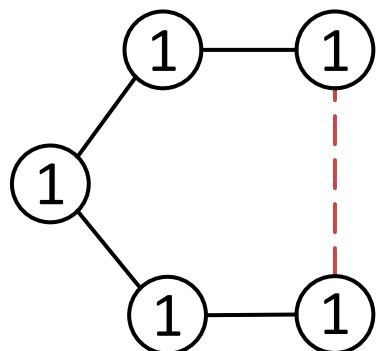
0	1	1	0	0
1	0	0	1	0
1	0	0	1	1
0	1	1	0	1
0	0	1	1	0
0	0	0	1	1
0	0	1	0	1
0	1	0	1	1
1	0	1	0	0
1	1	1	0	0

**Same graph**  
**Very different representations!**

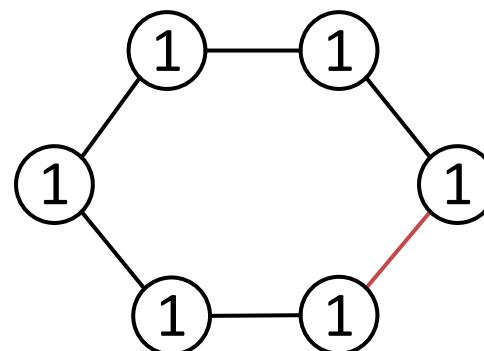
# Why is it Hard?

- **Complex dependencies:**
  - Edge formation has long-range dependencies

**Example: Generate a ring graph on 6 nodes:**



Shouldn't  
have edge!



Should  
have edge!

Existence of an edge may depend on the entire graph!

# Outline of Today's Lecture

- 1. Problem of Graph Generation**
- 2. ML Basics for Graph Generation**
- 3. GraphRNN**
- 4. Applications and Open Questions**



# **Machine Learning for Graph Generation**

# Graph Generative Models

- **Given:** Graphs sampled from  $p_{data}(G)$
- **Goal:**
  - Learn the distribution  $p_{model}(G)$
  - Sample from  $p_{model}(G)$

$p_{data}(G)$



Learn &  
Sample



$p_{model}(G)$



# Recap: Generative Models

## Setup:

- Assume we want to learn a generative model from a set of data points (i.e., graphs)  $\{\mathbf{x}_i\}$ 
  - $p_{data}(\mathbf{x})$  is the **data distribution**, which is never known to us, but we have sampled  $\mathbf{x}_i \sim p_{data}(\mathbf{x})$
  - $p_{model}(\mathbf{x}; \theta)$  is the **model**, parametrized by  $\theta$ , that we use to approximate  $p_{data}(\mathbf{x})$
- **Goal:**
  - (1) Make  $p_{model}(\mathbf{x}; \theta)$  close to  $p_{data}(\mathbf{x})$
  - (2) Make sure we can sample from  $p_{model}(\mathbf{x}; \theta)$ 
    - We need to generate examples (graphs) from  $p_{model}(\mathbf{x}; \theta)$

# Recap: Generative Models

## (1) Make $p_{model}(x; \theta)$ close to $p_{data}(x)$

- Key Principle: **Maximum Likelihood**
- Fundamental approach to modeling distributions

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{data}} \log p_{model}(x \mid \theta)$$

- Find parameters  $\theta^*$ , such that for observed data points  $x_i \sim p_{data}$ ,  $\sum_i \log p_{model}(x_i; \theta^*)$  has the highest value, among all possible choices of  $\theta$ 
  - That is, find the model that is most likely to have generated the observed data  $x$

# Recap: Generative Models

## (2) Sample from $p_{model}(x; \theta)$

- **Goal:** Sample from a complex distribution
- The most common approach:
  - (1) Sample from a simple noise distribution

$$\mathbf{z}_i \sim N(0, 1)$$

- (2) Transform the noise  $z_i$  via  $f(\cdot)$

$$x_i = f(\mathbf{z}_i; \theta)$$

Then  $x_i$  follows a complex distribution

- **Q:** How to design  $f(\cdot)$ ?
- **A:** Use Deep Neural Networks, and train it using the data we have!

# Recap: Generative Models

## Taxonomy of Deep Generative Models

This lecture:  
Auto-regressive  
models:

### Tractable density

- Fully visible belief nets
- NADE
- MADE
- PixelRNN
- Change of variables models (nonlinear ICA)

### Explicit density

### Implicit density

### Approximate density

#### Variational

#### Markov Chain

Variational autoencoder      Boltzmann machine

...

### Maximum Likelihood

### Direct GAN

### Markov Chain GSN

An autoregressive (AR) model predicts future behavior based on past behavior.

# Recap: Generative Models

## Auto-regressive models

- $p_{model}(x; \theta)$  is used for **both density estimation and sampling** (from the probability density)
  - (other models like Variational Auto Encoders (VAEs), Generative Adversarial Nets (GANs) have 2 or more models, each playing one of the roles)
- **Apply chain rule:** Joint distribution is a product of conditional distributions:

$$p_{model}(x; \theta) = \prod_{t=1}^n p_{model}(x_t | x_1, \dots, x_{t-1}; \theta)$$

- E.g.,  $x$  is a vector,  $x_t$  is the  $t$ -th dimension;  
 $x$  is a sentence,  $x_t$  is the  $t$ -th word.
- **In our case:**  $x_t$  will be the  $t$ -th action (add node, add edge)

# Outline of Today's Lecture

- 1. Problem of Graph Generation**
- 2. ML Basics for Graph Generation**
- 3. GraphRNN** 
- 4. Applications and Open Questions**

# GraphRNN: Generating Realistic Graphs

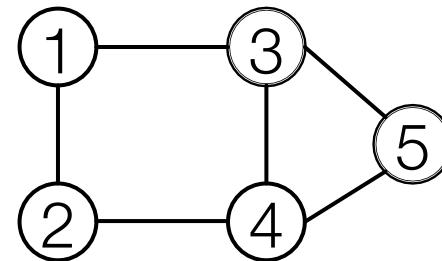
[GraphRNN: Generating Realistic Graphs with Deep Auto-regressive Models](#). J. You, R. Ying, X. Ren, W. L. Hamilton, J. Leskovec. *International Conference on Machine Learning (ICML)*, 2018.

# GraphRNN Idea

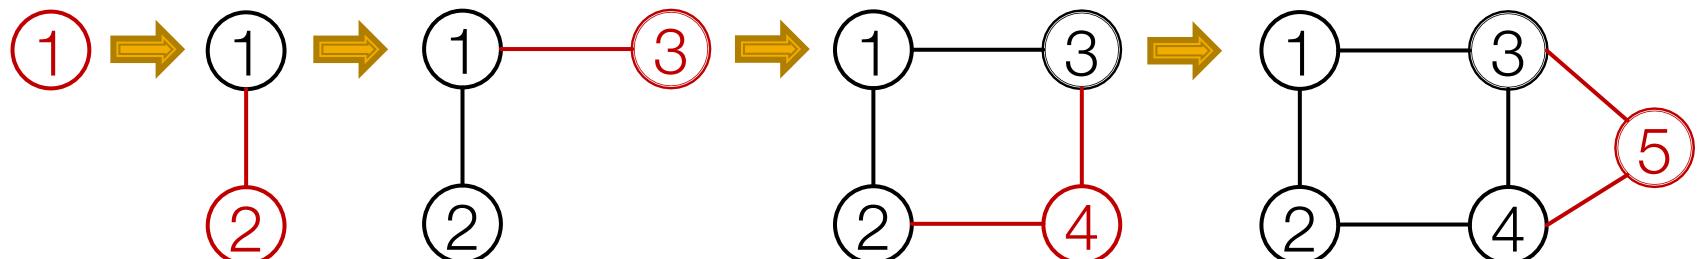
[You et al., ICML 2018]

**Generating graphs via sequentially adding nodes and edges**

Graph  $G$



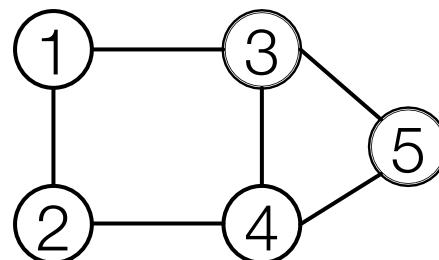
Generation process  $S^\pi$



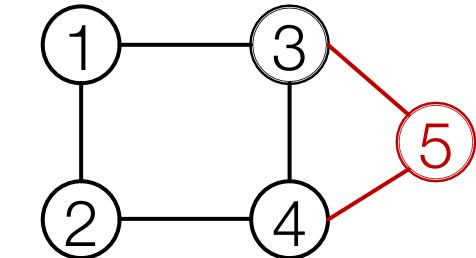
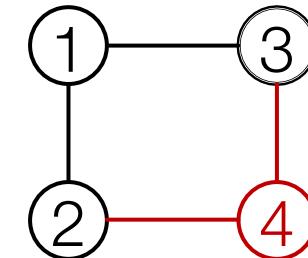
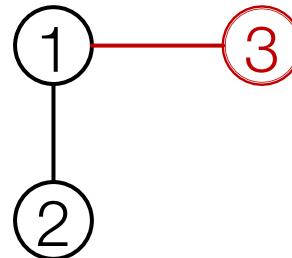
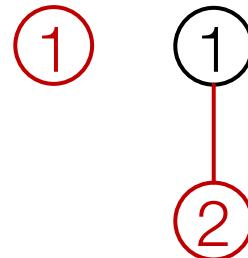
# Model Graphs as Sequences

Graph  $G$  with node ordering  $\pi$  can be uniquely mapped into a sequence of node and edge additions  $S^\pi$

Graph  $G$  with  
node ordering  $\pi$ :



Sequence  $S^\pi$ :



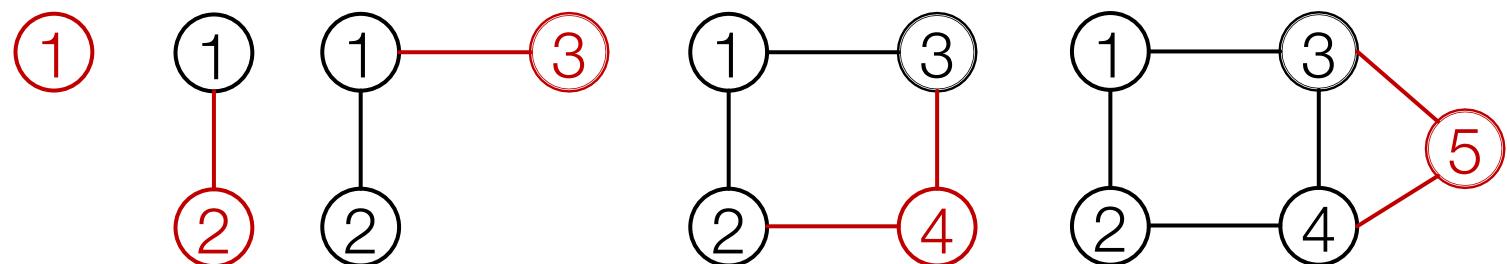
$$S^\pi = ( S_1^\pi, S_2^\pi, S_3^\pi, S_4^\pi, S_5^\pi )$$

# Model Graphs as Sequences

The sequence  $S^\pi$  has two levels

( $S$  is a sequence of sequences):

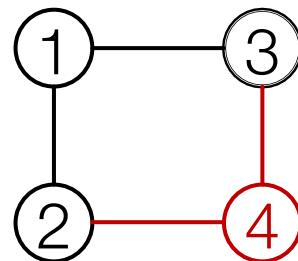
- **Node-level:** add nodes, one at a time
  - **Edge-level:** add edges between existing nodes  
  - **Node-level:** At each step, a **new node is added**



# Model Graphs as Sequences

The sequence  $S^\pi$  has **two levels**:

- Each **Node-level** step is an **edge-level** sequence
- **Edge-level:** At each step, add a new edge



$$S_4^\pi$$

$$S_4^\pi = ( S_{4,1}^\pi , \quad S_{4,2}^\pi , \quad S_{4,3}^\pi )$$

“Not connect 4, 1”   “Connect 4, 2”   “Connect 4, 3”

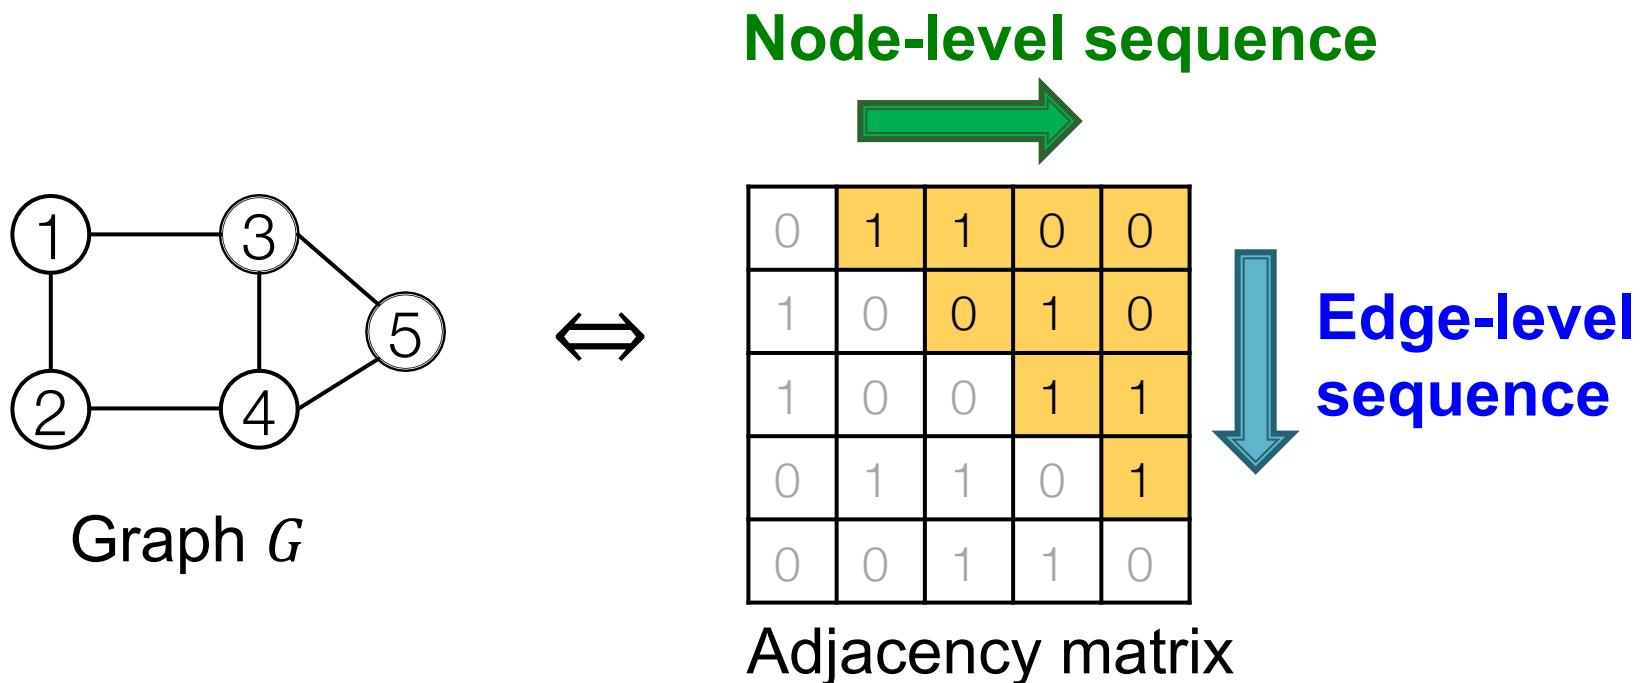
0

1

1

# Model Graphs as Sequences

- Summary: A graph + a node ordering = A sequence of sequences!
- Node ordering is randomly selected (we will come back to this)



# Model Graphs as Sequences

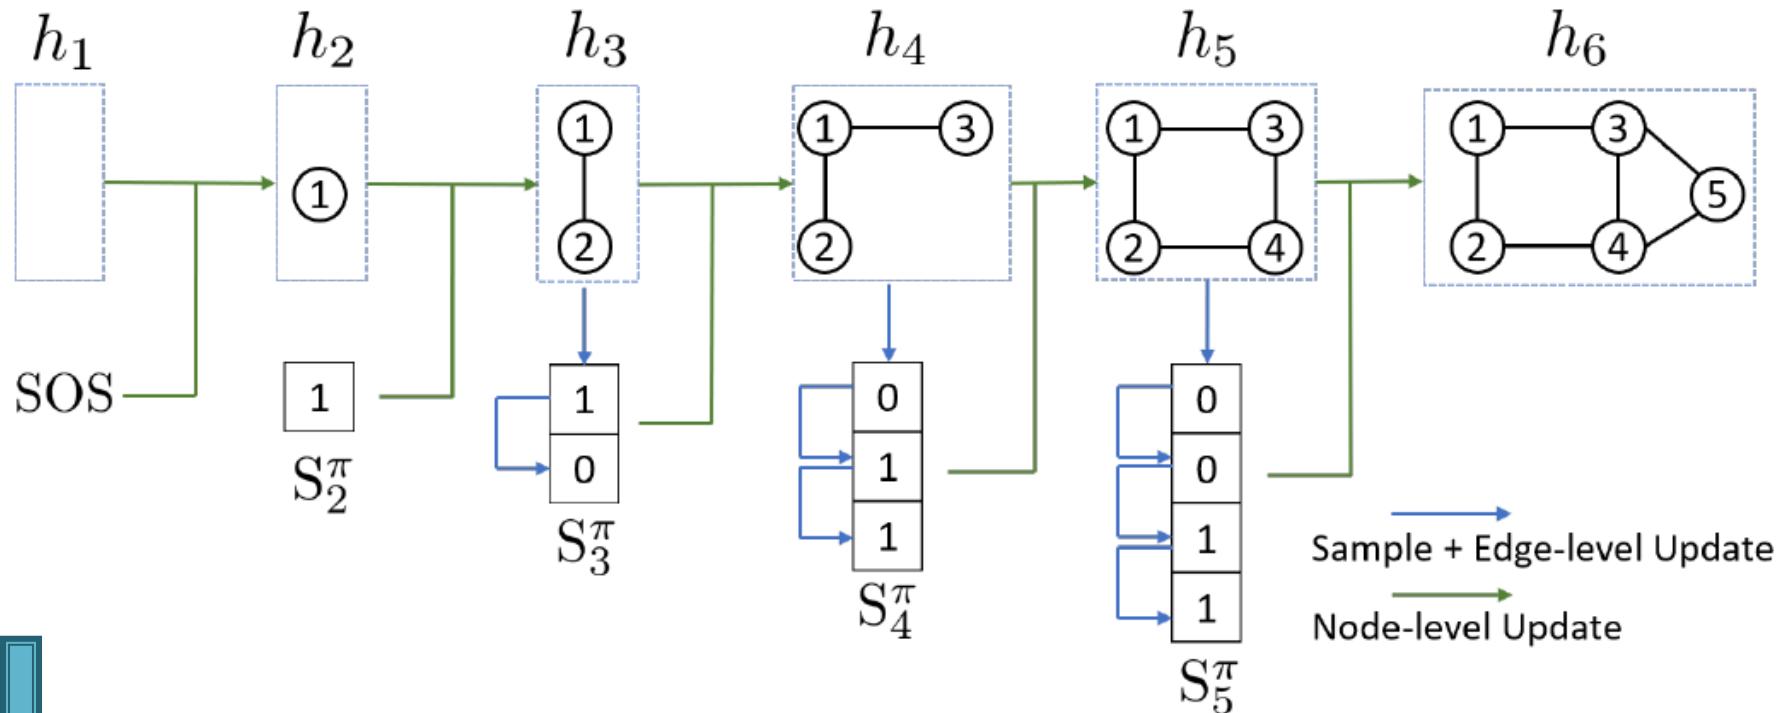
- We have transformed graph generation problem into a sequence generation problem
- Need to model two processes:
  - Generate a state for a new node  
(Node-level sequence)
  - Generate edges for the new node based on its state  
(Edge-level sequence)
- Approach: Use RNN to model these processes!

# GraphRNN: Two levels of RNN

- GraphRNN has a **node-level RNN** and an **edge-level RNN**
- Relationship between the two RNNs:
  - Node-level RNN generates the initial state for edge-level RNN
  - Edge-level RNN generates edges for the new node, then updates node-level RNN state using generated results

# GraphRNN: Two levels of RNN

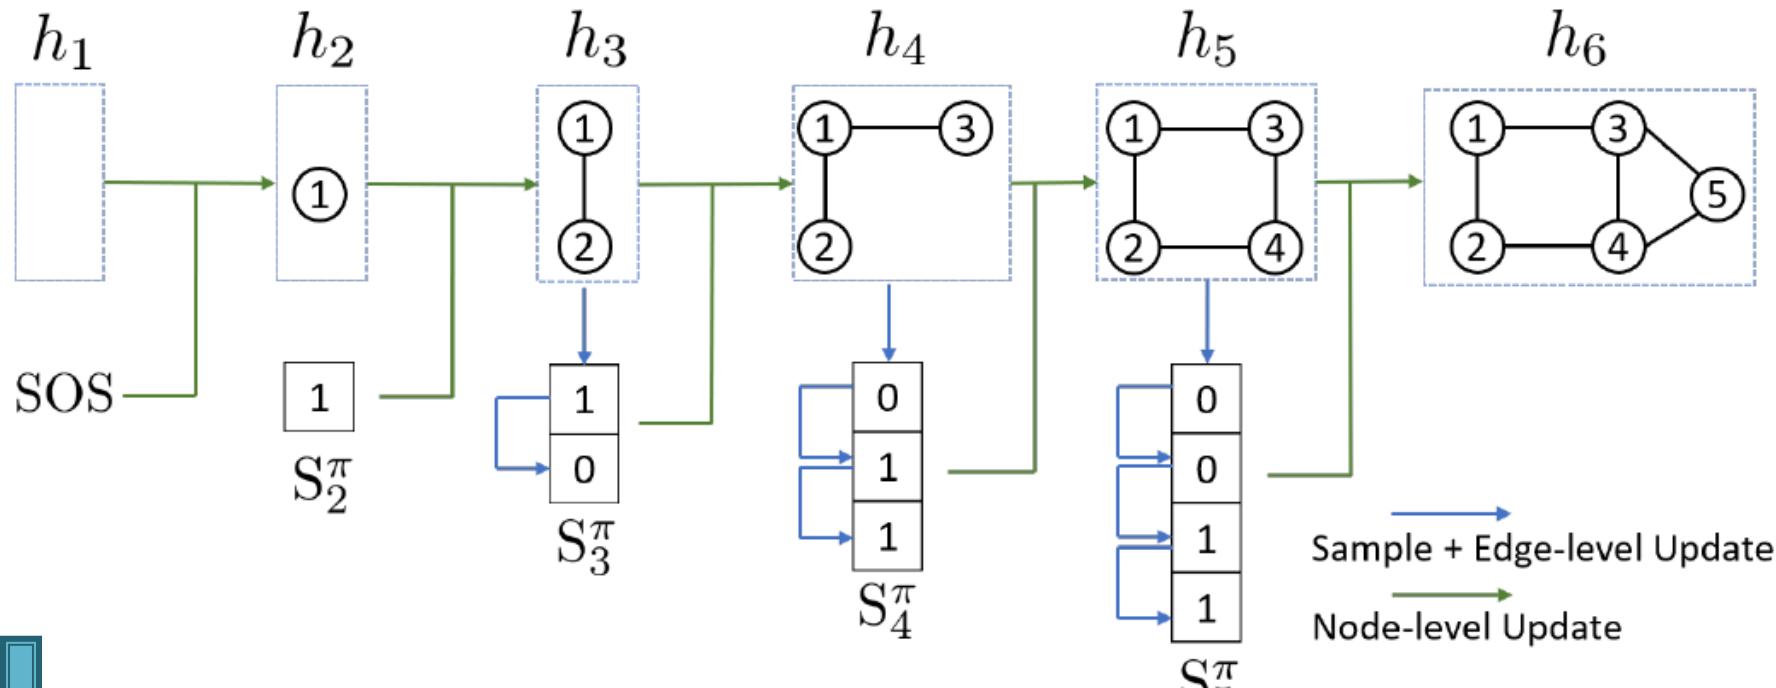
Node-level RNN generates the initial state for edge-level RNN



Edge-level RNN generates edges for the new node,  
then update node-level RNN state using generated results

# GraphRNN: Two levels of RNN

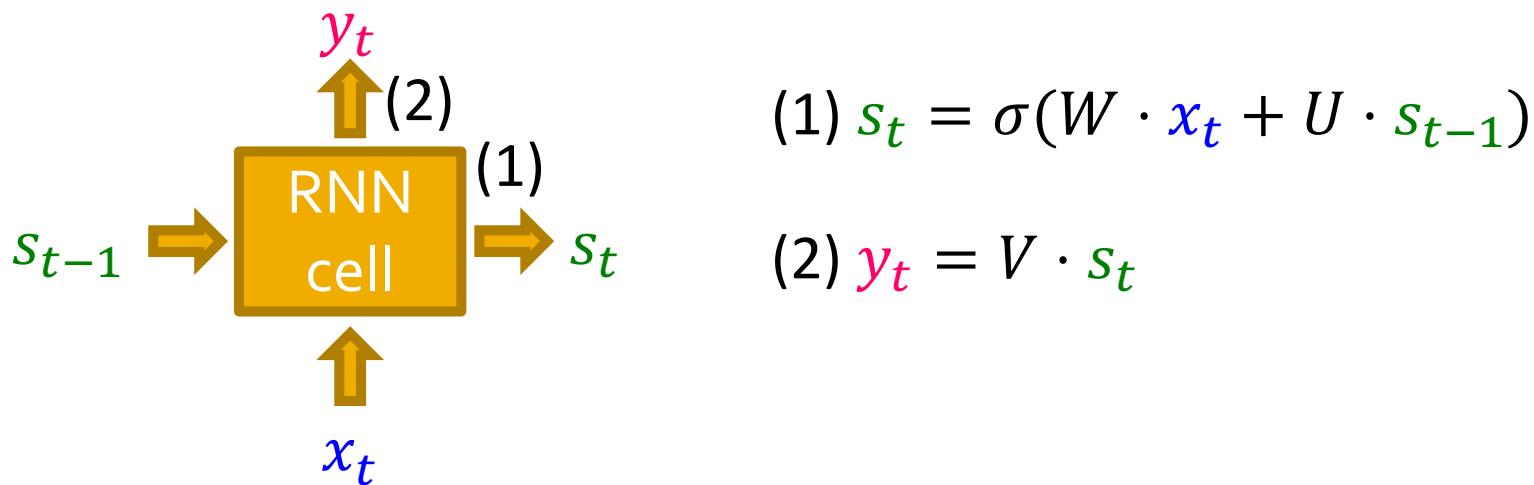
Node-level RNN generates the initial state for edge-level RNN



Next: How to generate a sequence with RNN?

# Recurrent Neural Networks

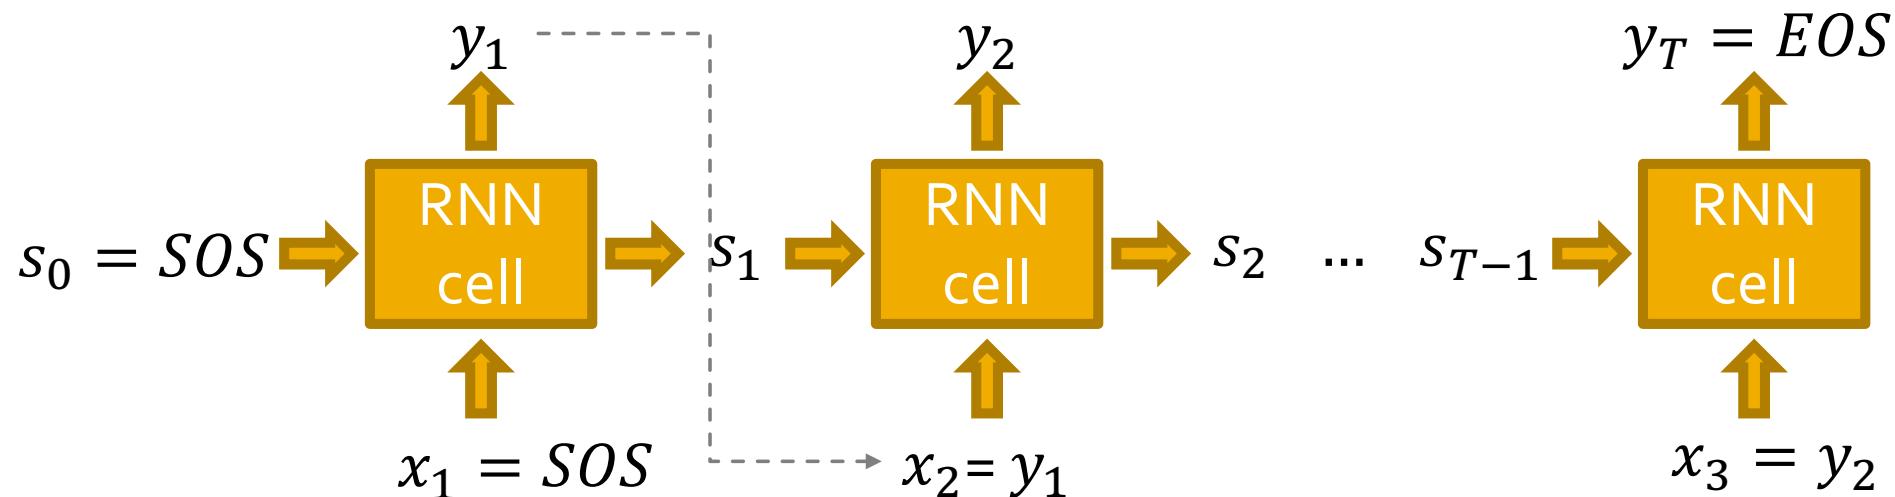
- $s_t$ : State of RNN after time  $t$
- $x_t$ : Input to RNN at time  $t$
- $y_t$ : Output of RNN at time  $t$
- $W, U, V$ : parameter matrices,  $\sigma(\cdot)$ : non-linearity



- More expressive cells: GRU, LSTM, etc.

# RNN for Sequence Generation

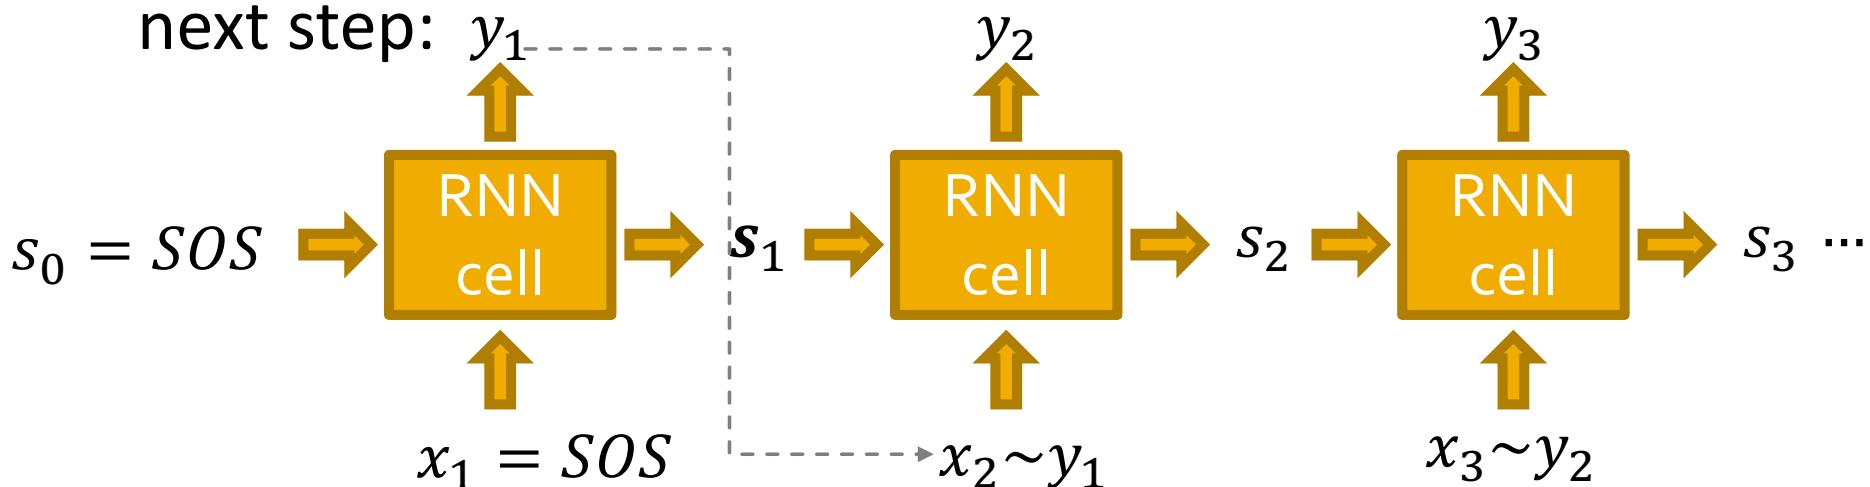
- **Q:** How to use RNN to generate sequences?
- **A:** Let  $x_{t+1} = y_t$ !
- **Q:** How to initialize  $s_0, x_1$ ? When to stop generation?
- **A:** Use start/end of sequence token (SOS, EOS)- e.g., zero vector



- This is good, but this model is deterministic

# RNN for Sequence Generation

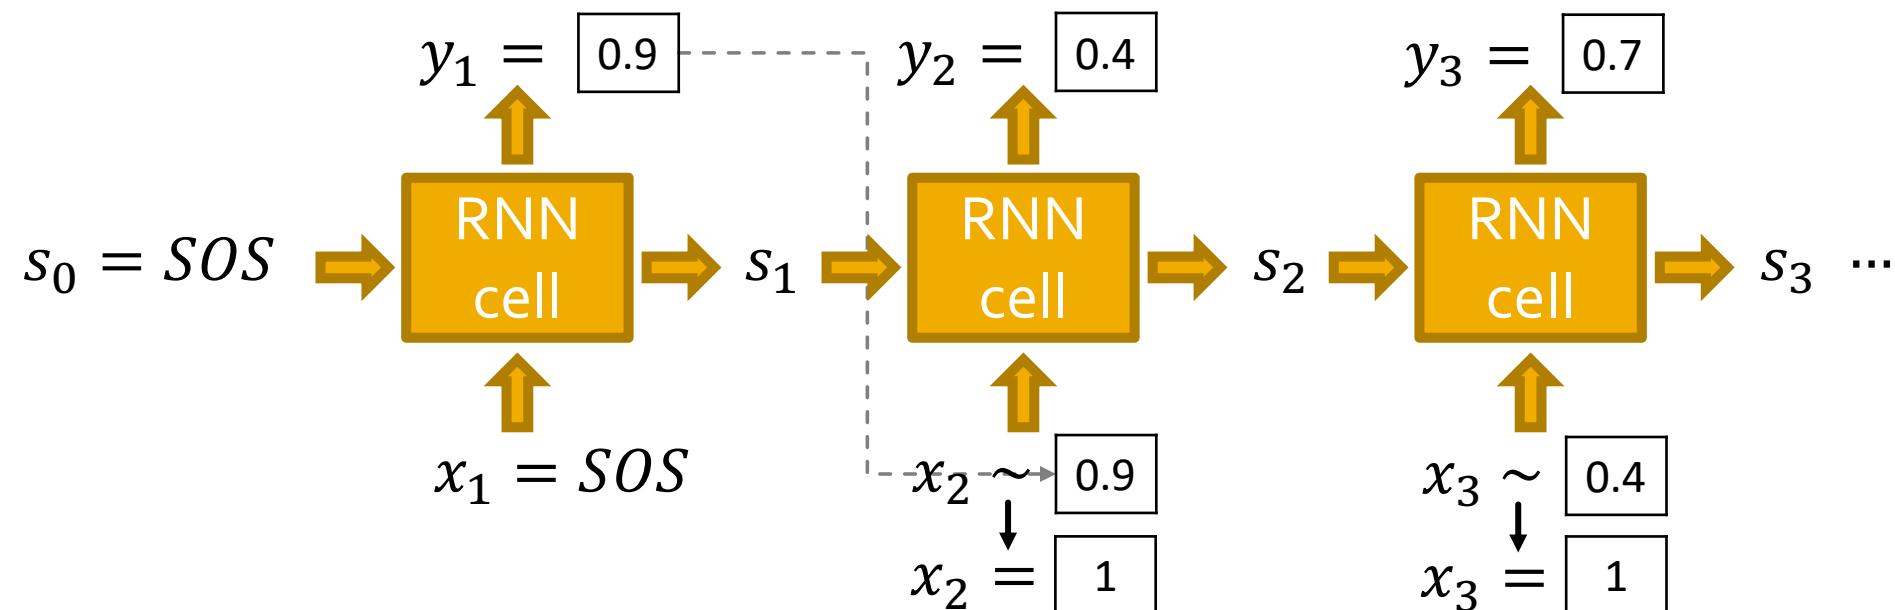
- Remember our goal: Use RNN to model  $\prod_{k=1}^n p_{model}(x_t | x_1, \dots, x_{t-1}; \theta)$
- Let  $y_t = p_{model}(x_t | x_1, \dots, x_{t-1}; \theta)$
- Then  $x_{t+1}$  is a sample from  $y_t$ :  $x_{t+1} \sim y_t$ 
  - Each step of RNN outputs a probability vector
  - We then sample from the vector, and feed sample to next step:



# RNN at Test Time

Suppose we already have trained the model

- $y_t$  follows Bernoulli distribution (choice of  $p_{model}$ )
- $\boxed{p}$  means value 1 has prob.  $p$ , value 0 has prob.  $1 - p$

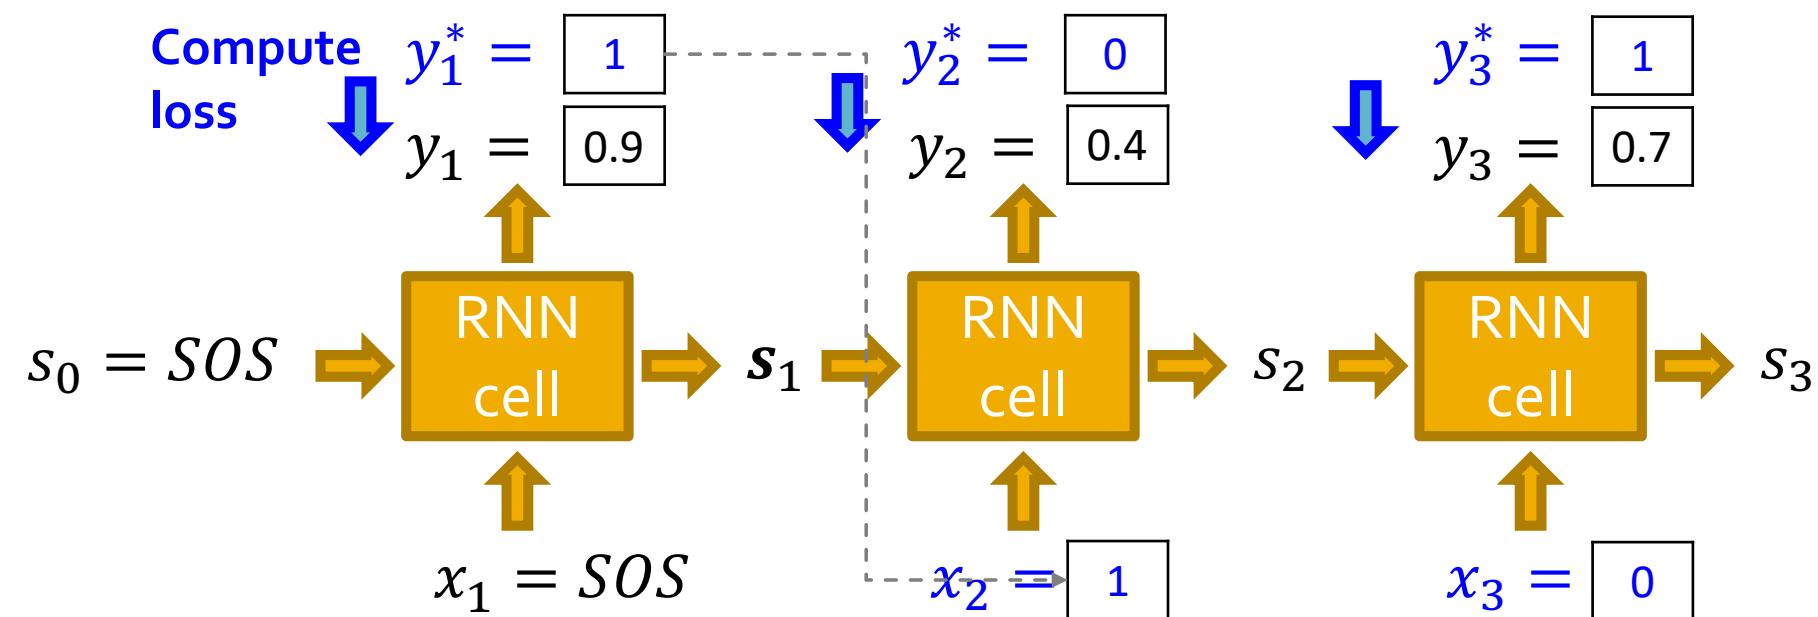


- Right now everything is generated by the model
- **How do we use training data  $x_1, x_2, \dots, x_n$ ?**

# RNN at Training Time

## Training the model

- We observe a sequence  $y^*$  of edges [1,0,...]
- Principle: Teacher Forcing** -- Replace input and output by the real sequence



# RNN at Training Time

- Loss  $L$  : **Binary cross entropy**
- Minimize:

$$L = -[y_1^* \log(y_1) + (1 - y_1^*) \log(1 - y_1)]$$

Compute  
loss

$\downarrow$

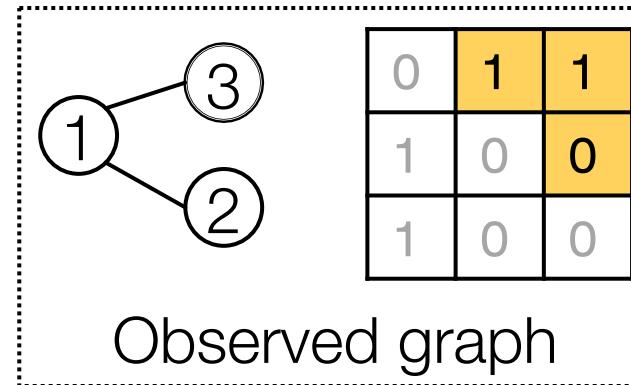
$$\begin{aligned} y_1^* &= \boxed{1} \\ y_1 &= \boxed{0.9} \end{aligned}$$

- If  $y_1^* = 1$ , we minimize  $-\log(y_1)$ , making  $y_1$  higher
- If  $y_1^* = 0$ , we minimize  $-\log(1 - y_1)$ , making  $y_1$  lower
- This way,  $y_1$  is **fitting** the data samples  $y_1^*$
- **Reminder:**  $y_1$  is computed by RNN, this loss will **adjust RNN parameters accordingly**, using back propagation!

# Put Things Together: Training

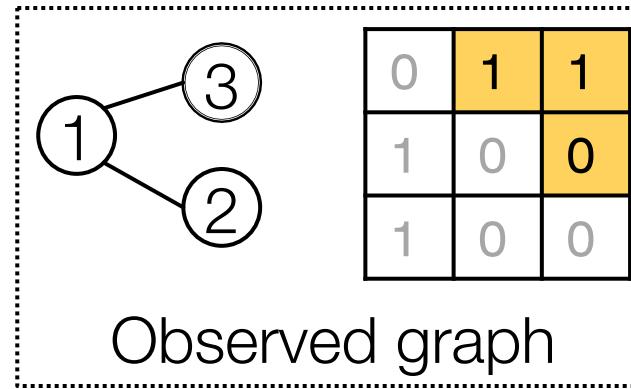
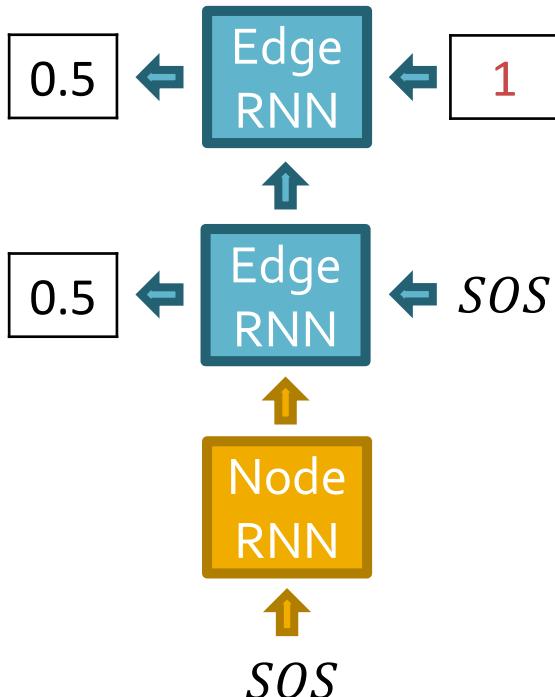
Assuming Node 1 is in the graph

Now adding Node 2



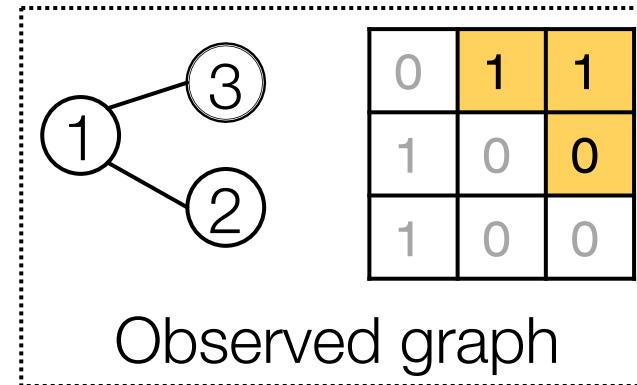
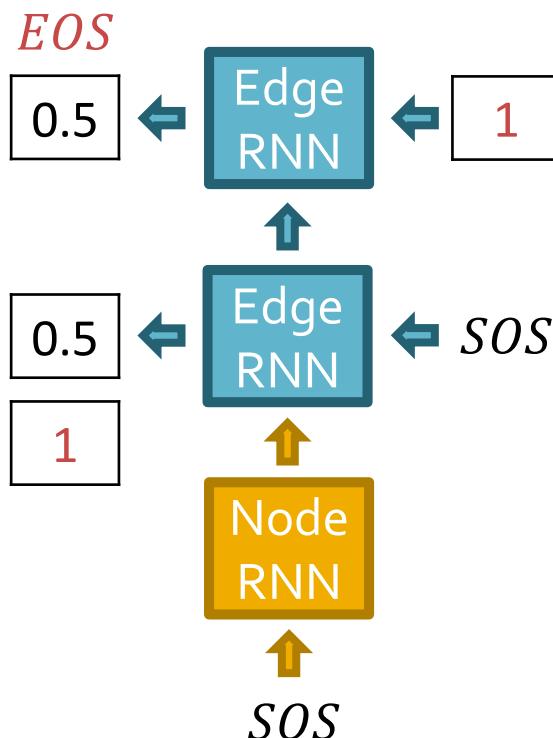
# Put Things Together: Training

Edge RNN predicts how  
Node 2 connects to Node 1



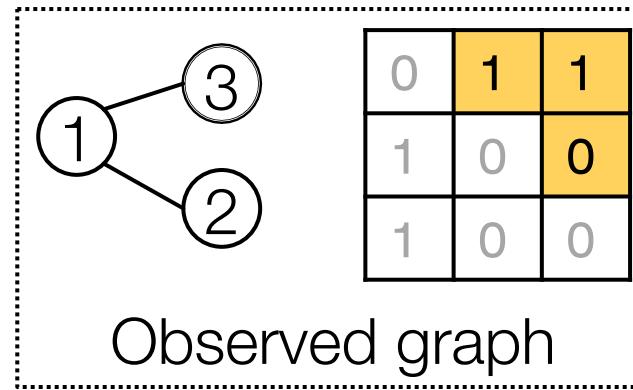
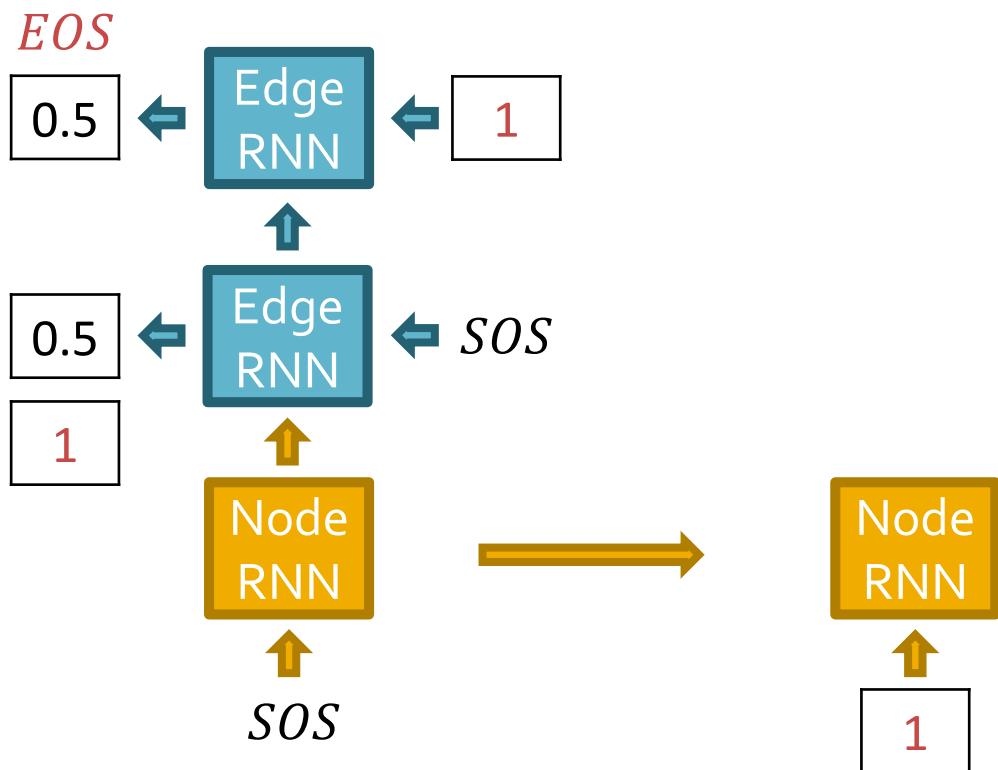
# Put Things Together: Training

Edge RNN gets supervisions from ground truth



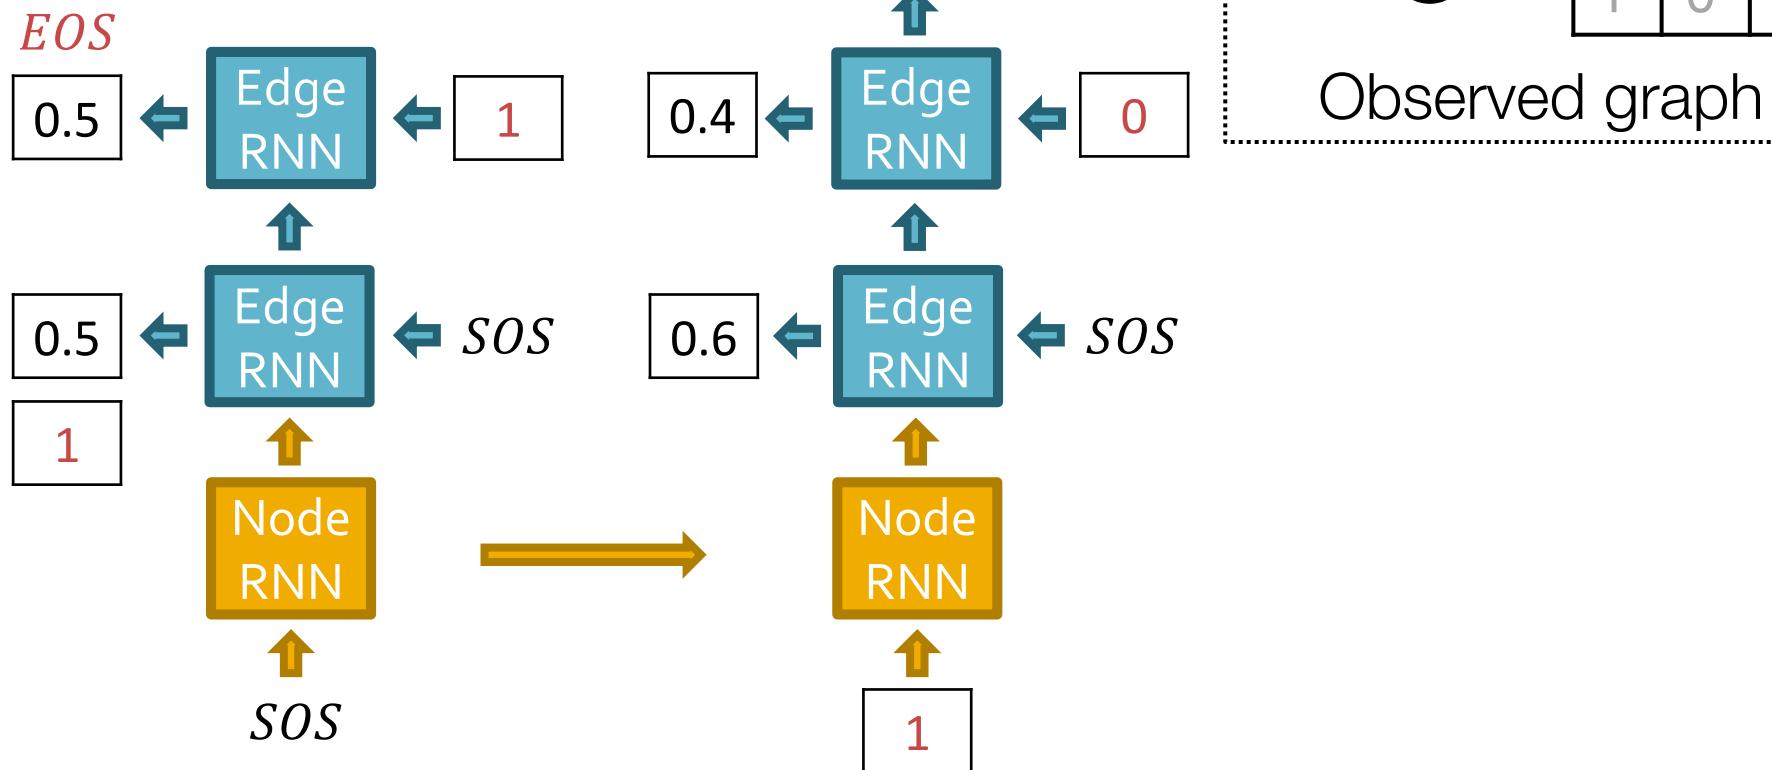
# Put Things Together: Training

New edges are used to update Node RNN



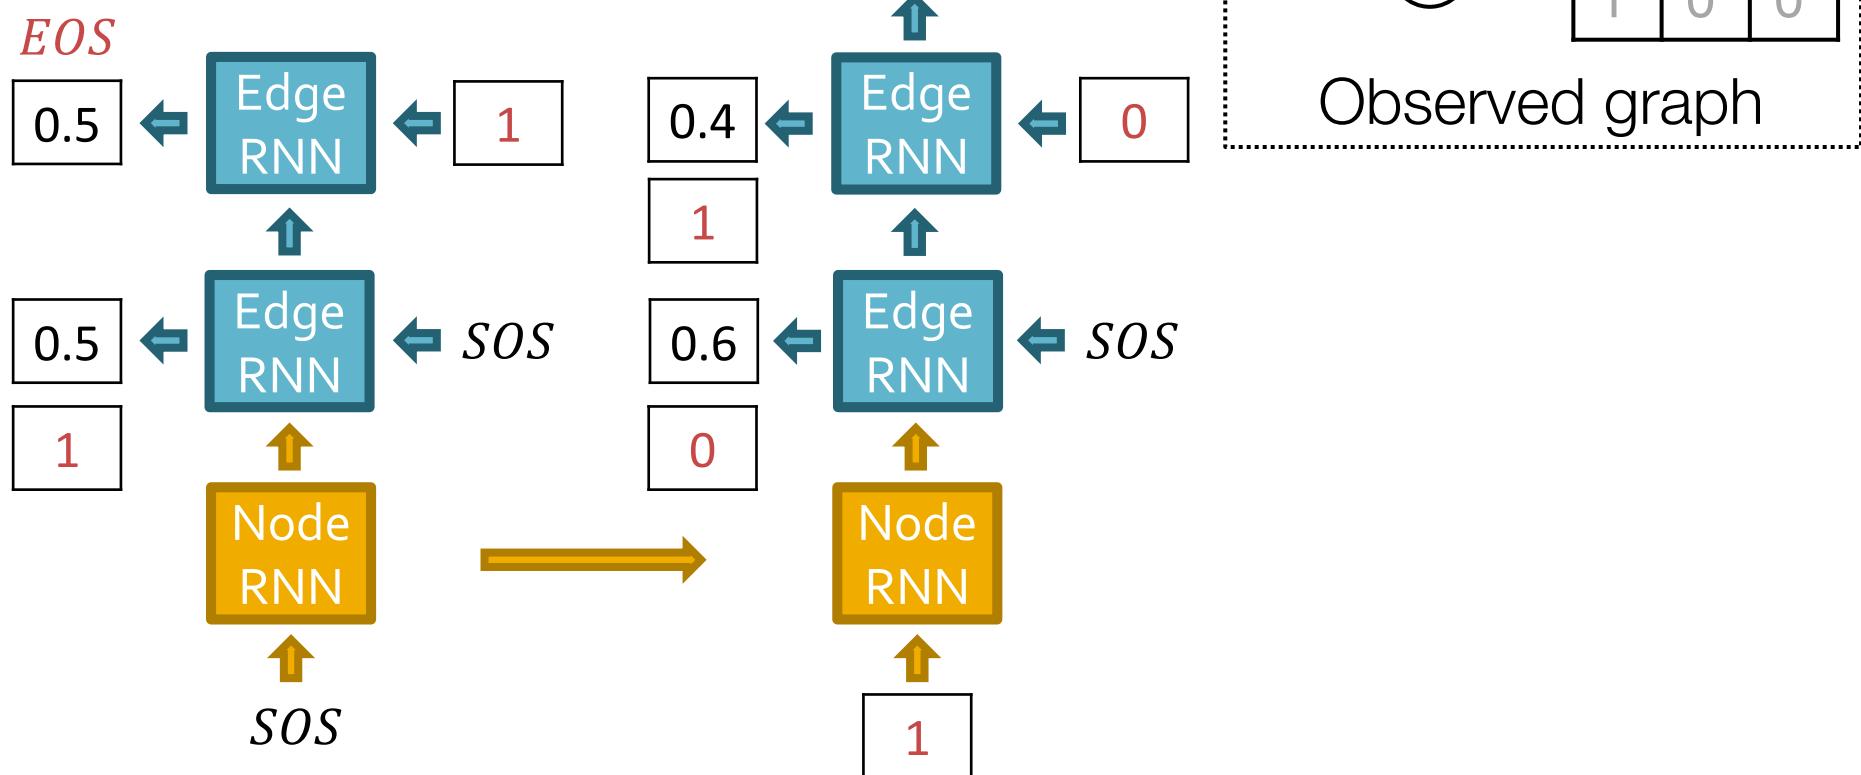
# Put Things Together: Training

Edge RNN predicts  
how Node 3 connects  
to Node 2



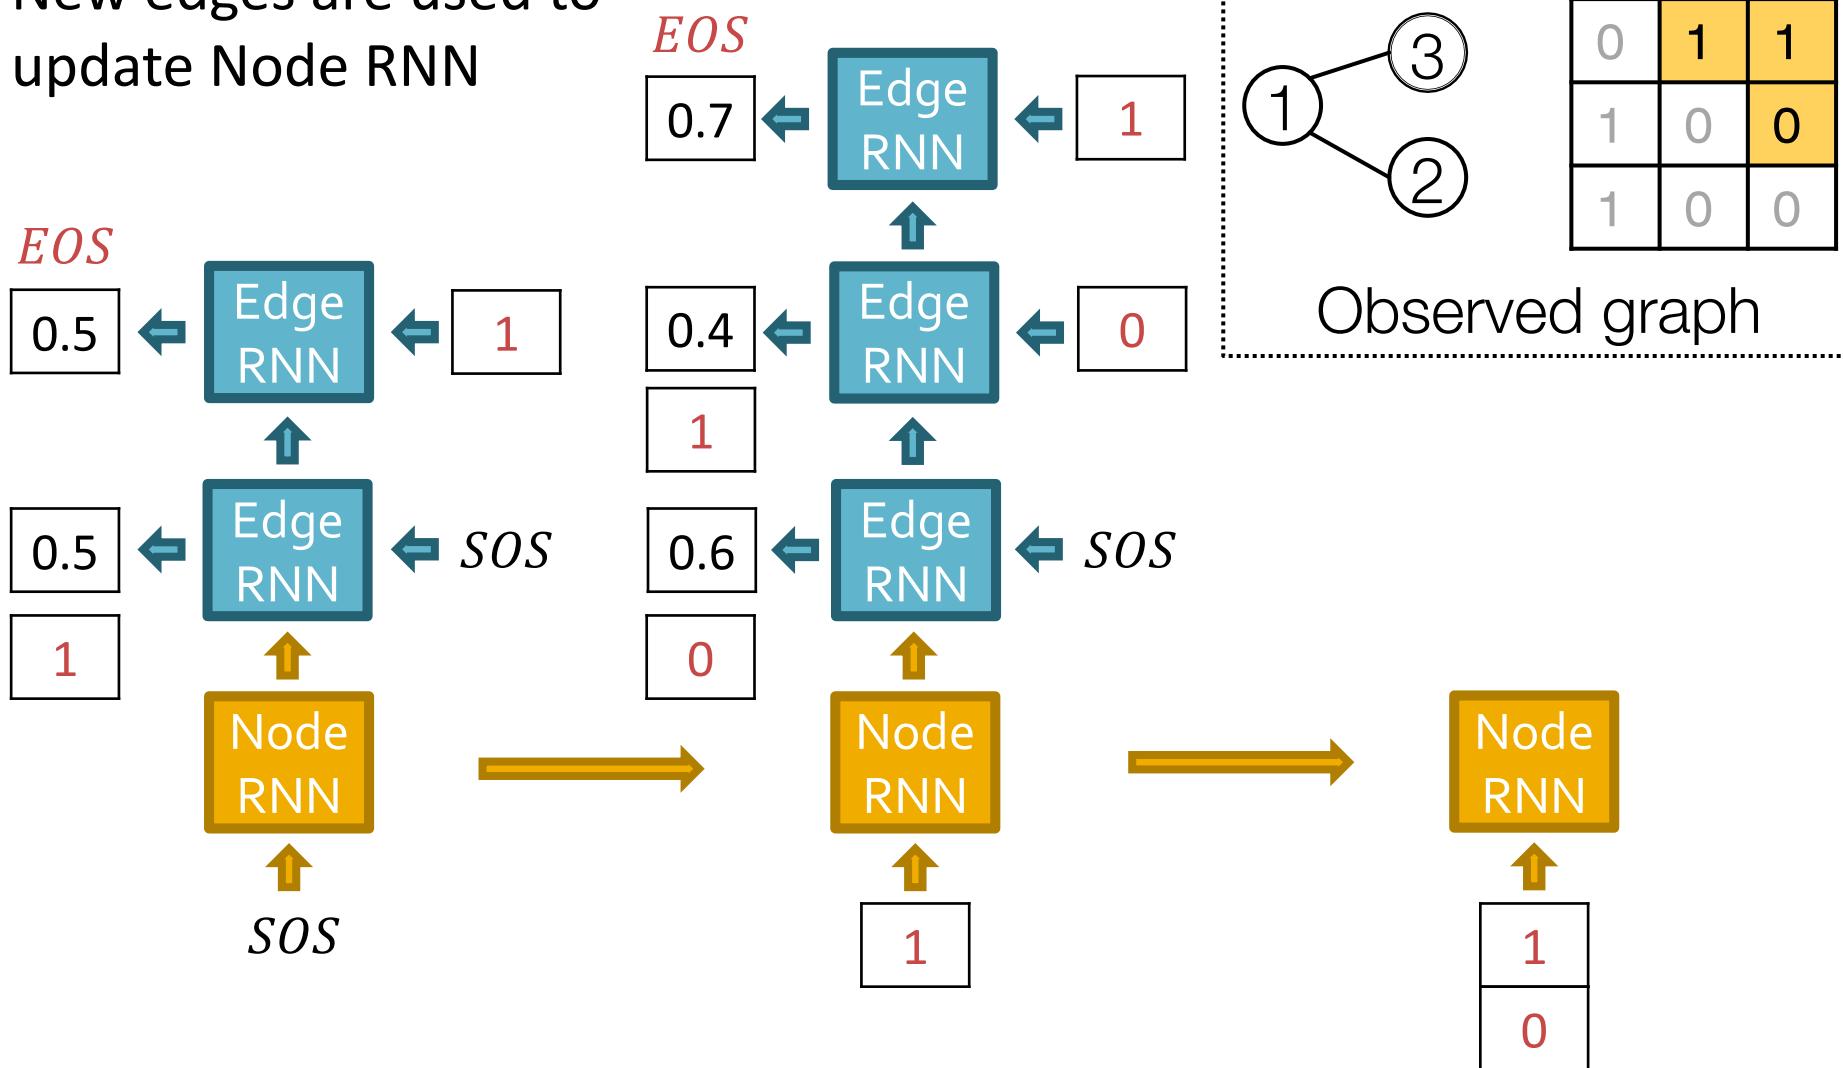
# Put Things Together: Training

Edge RNN gets supervisions from ground truth



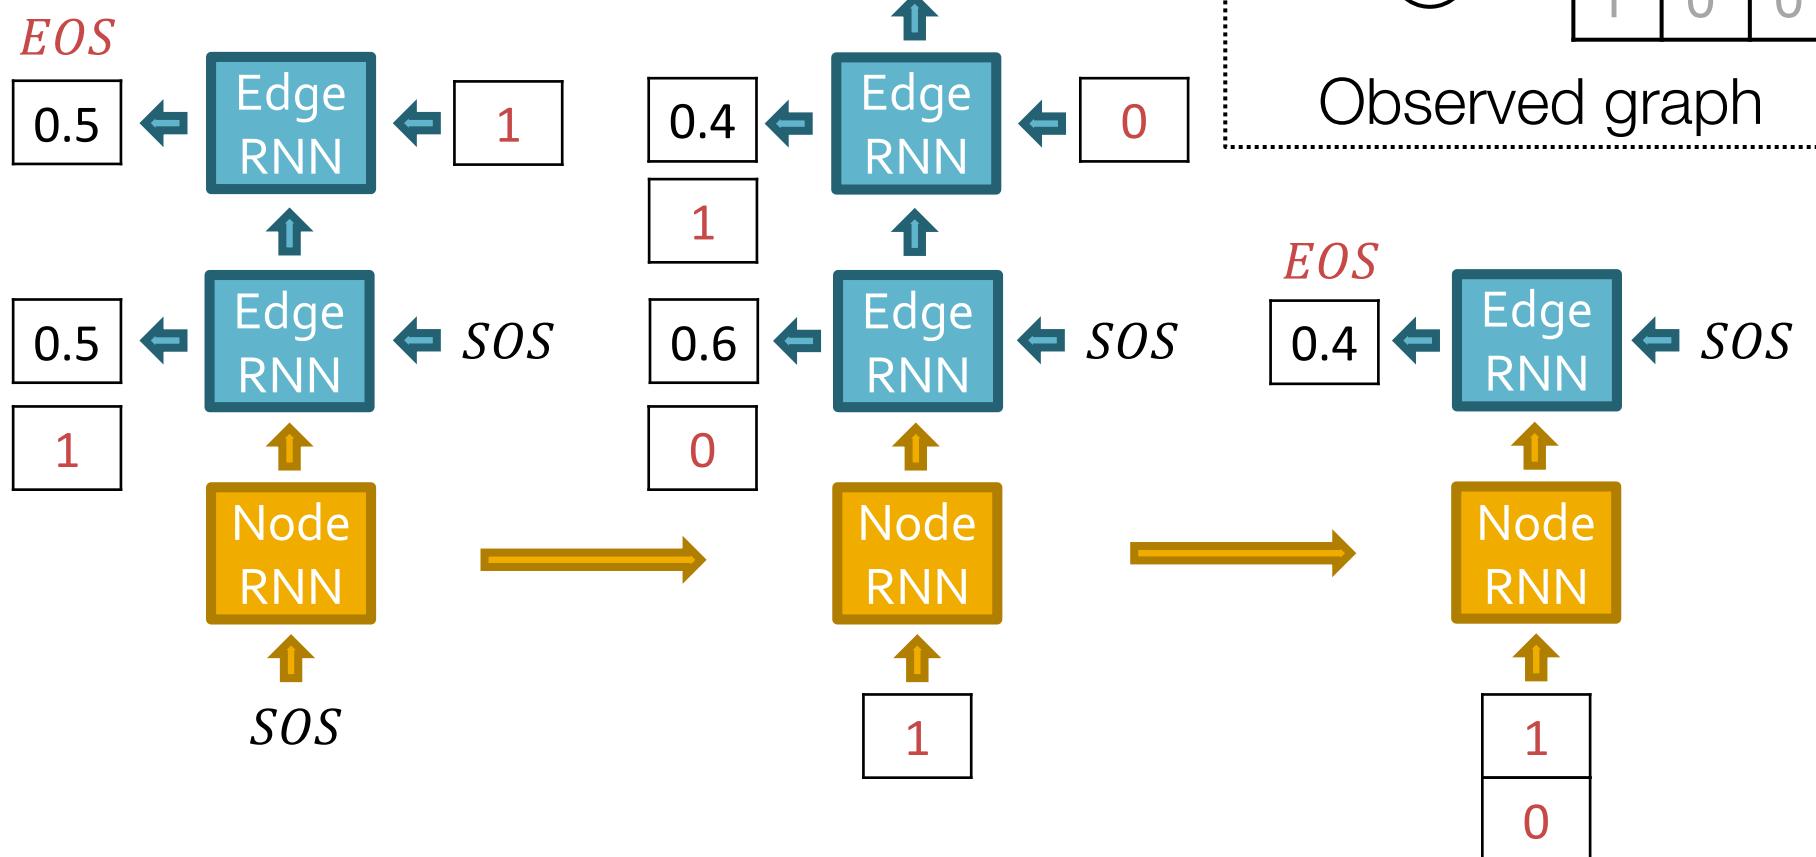
# Put Things Together: Training

New edges are used to update Node RNN



# Put Things Together: Training

Node 4 doesn't connect to any nodes, stop generation

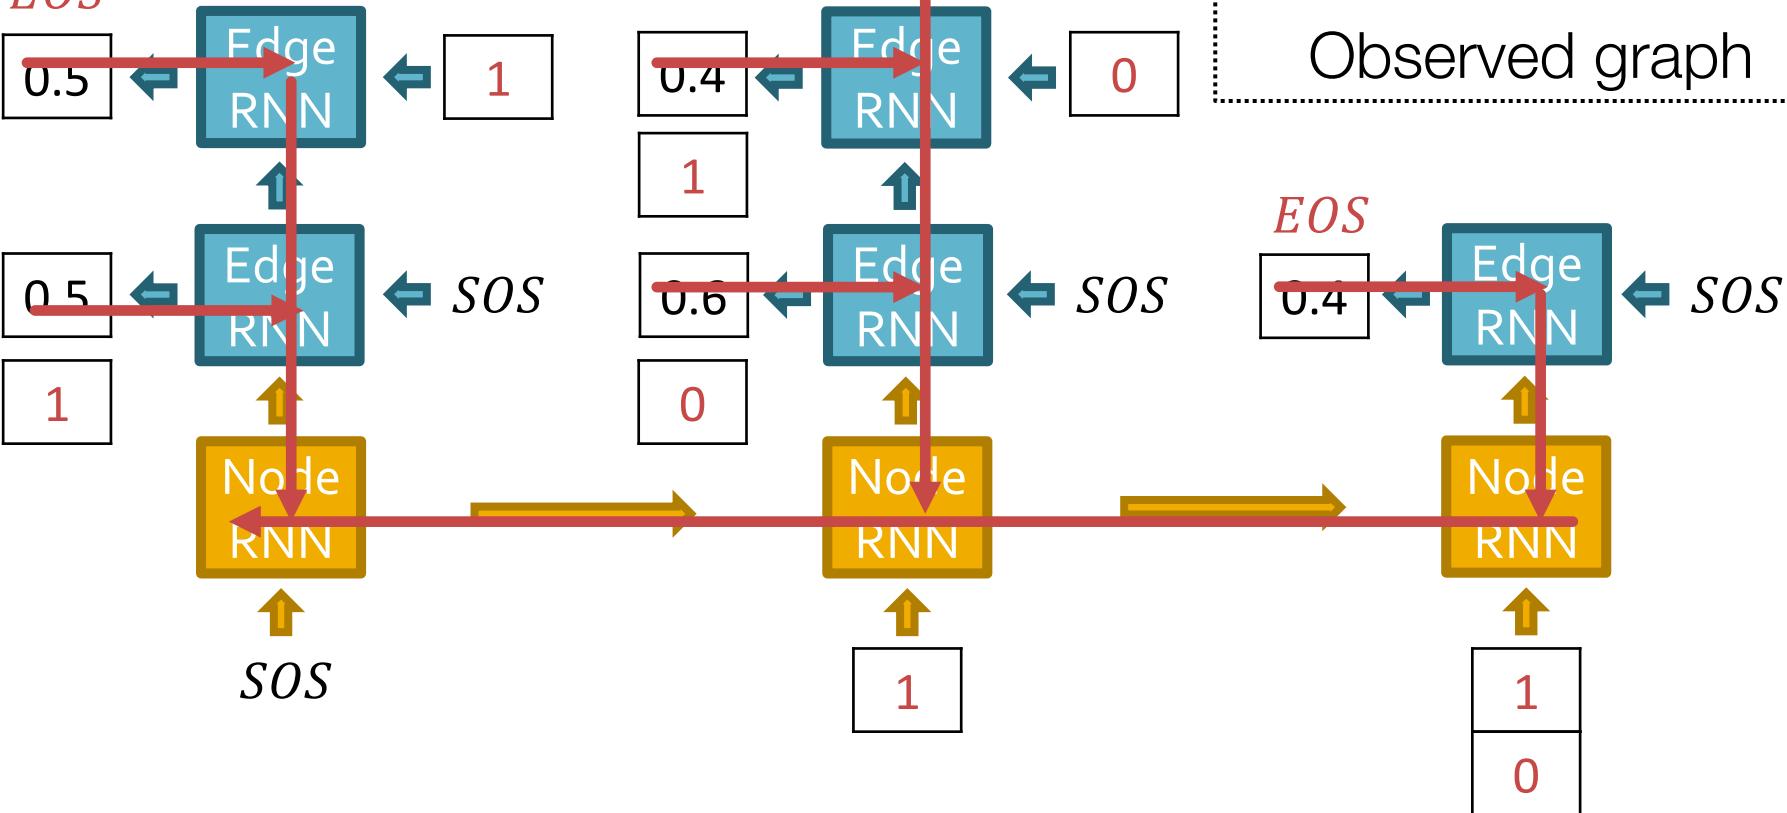


# Put Things Together: Training

Backprop through time:

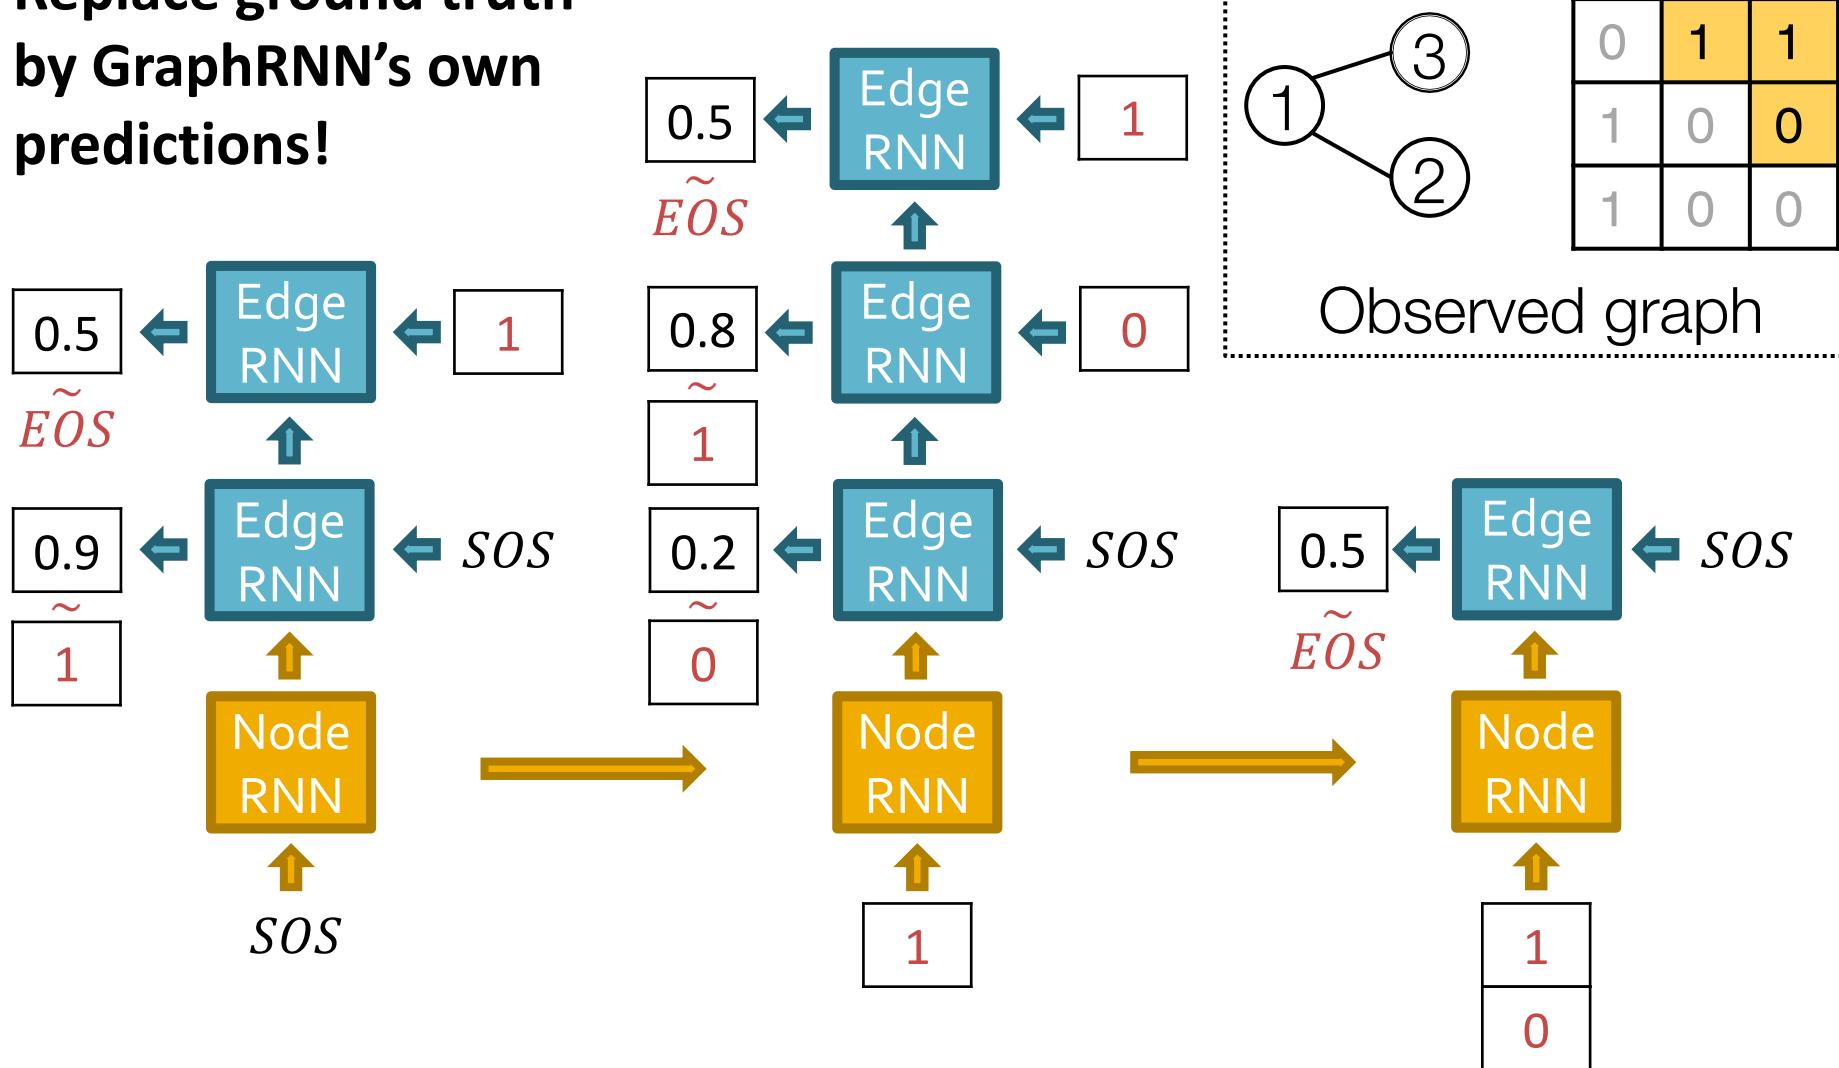
All gradients are  
accumulated across time  
steps

*EOS*



# Put Things Together: Test

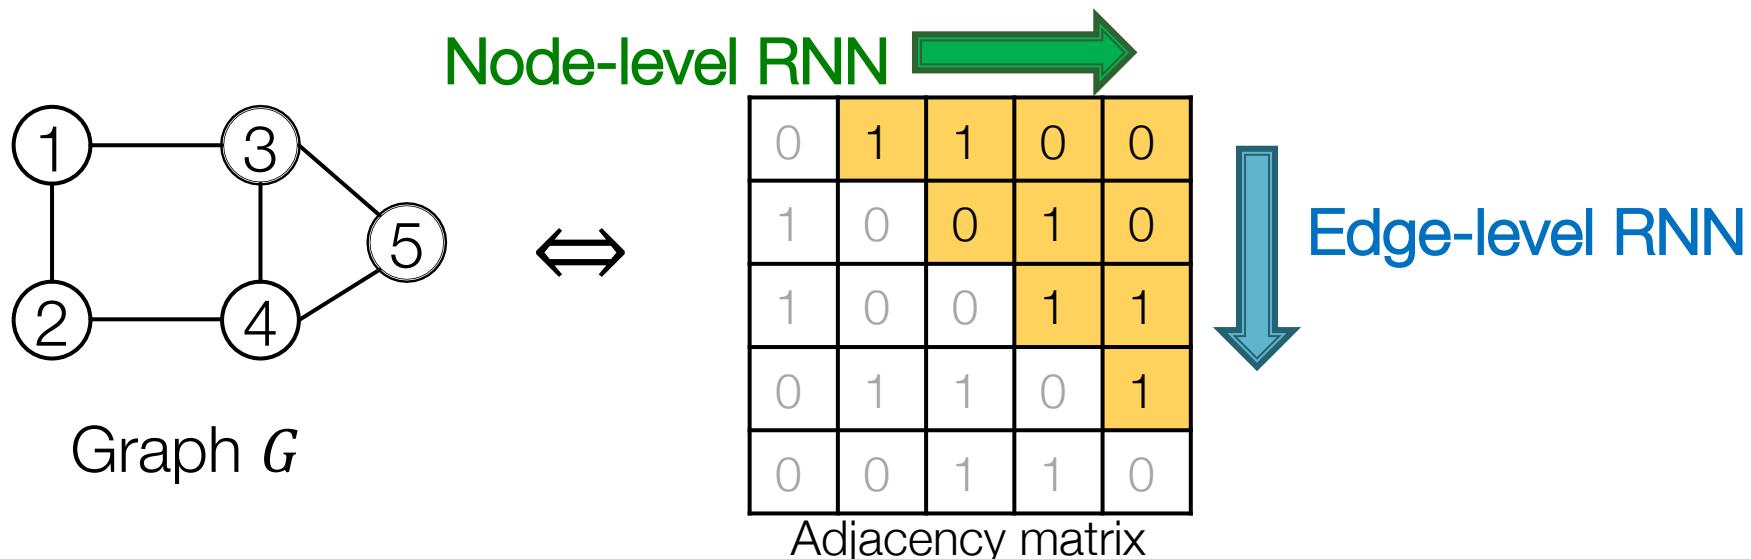
Replace ground truth  
by GraphRNN's own  
predictions!



# GraphRNN: Two levels of RNN

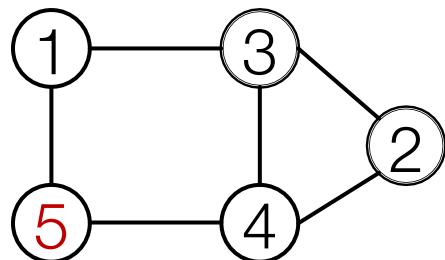
## Quick Summary of GraphRNN:

- Generate a graph by generating a two level sequence
- Use RNN to generate the sequences
- **Next:** Making GraphRNN tractable, proper evaluation



# Issue: Tractability

- Any node can connect to any prior node
- Too many steps for edge generation
  - Need to generate full adjacency matrix
  - Complex too-long edge dependencies



Random node ordering:

Node 5 may connect to any/all previous nodes

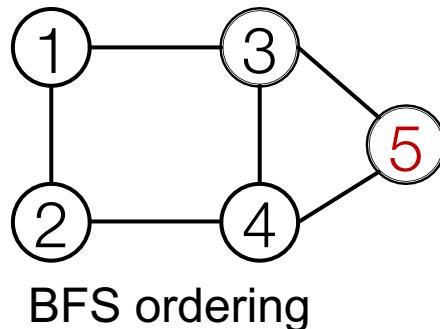
**“Recipe” to generate the left graph:**

- Add node 1
- Add node 2
- Add node 3
- Connect 3 with 1 and 2
- Add node 4
- ...

How do we limit this complexity?

# Solution: Tractability via BFS

## ■ Breadth-First Search node ordering



“Recipe” to generate the left graph:

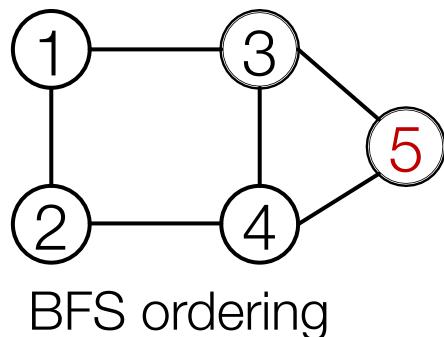
- Add node 1
- Add node 2
- Connect 2 with 1
- Add node 3
- Connect 3 with 1
- Add node 4
- Connect 4 with 2 and 3

## ■ BFS node ordering:

- Since Node 4 doesn't connect to Node 1
- We know all Node 1's neighbors have already been traversed
- Therefore, Node 5 and the following nodes will never connect to node 1
- We only need memory of 2 “steps” rather than  $n - 1$  steps

# Solution: Tractability via BFS

## ■ Breadth-First Search node ordering



BFS node ordering: Node 5 will never connect to node 1  
(only need memory of 2 “steps” rather than  $n - 1$  steps)

## ■ Benefits:

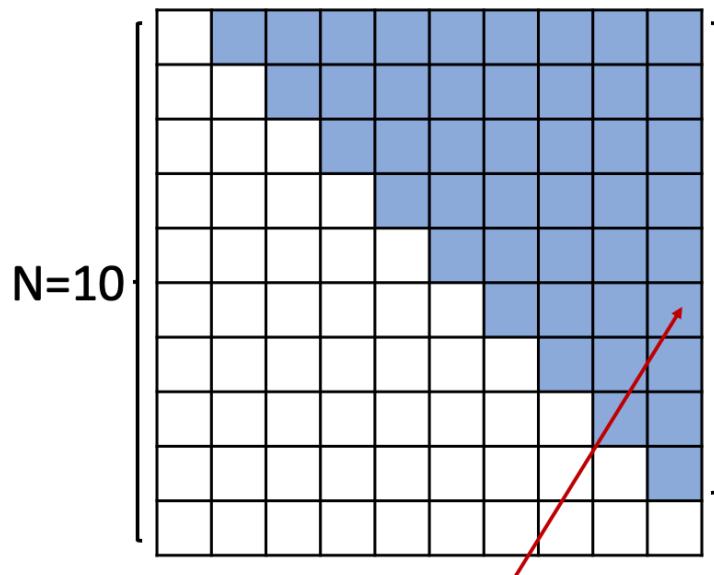
- Reduce possible node orderings
  - From  $O(n!)$  to number of distinct BFS orderings
- Reduce steps for edge generation
  - Reducing number of previous nodes to look at

# Solution: Tractability via BFS

- BFS reduces the number of steps for edge generation

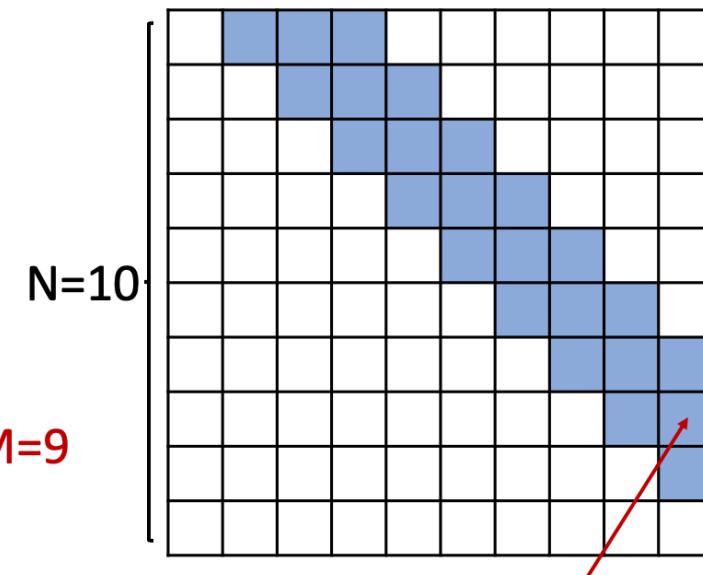
Adjacency matrices

Without BFS ordering



Connectivity with  
All Previous nodes

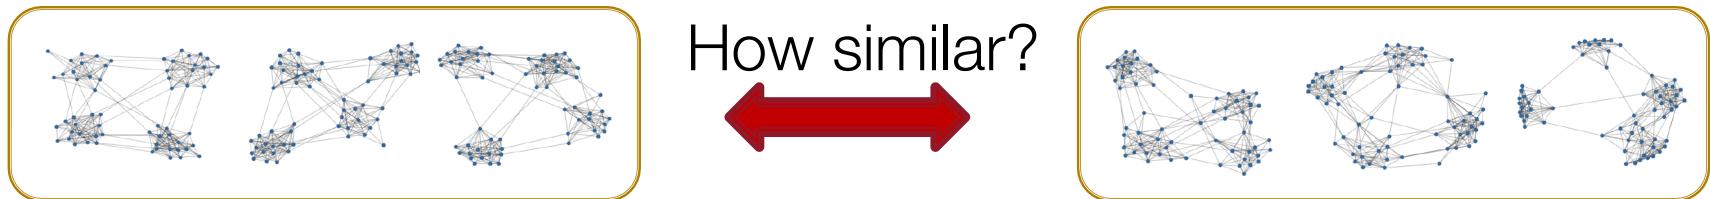
With BFS ordering



Connectivity only with  
nodes in the BFS frontier

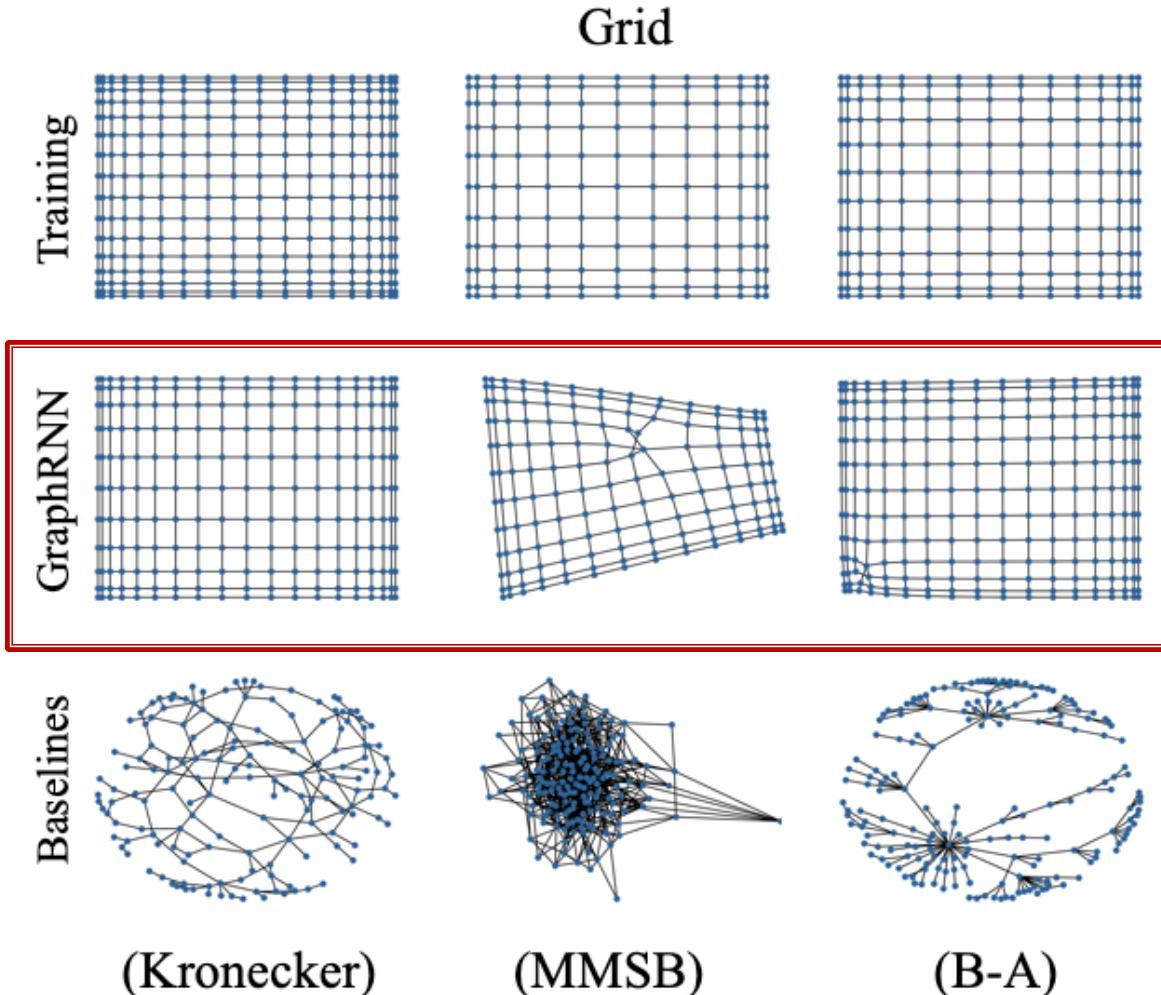
# Evaluating Generated Graphs

- Task: Compare two sets of graphs



- **Goal:** Define similarity metrics for graphs
- **Challenge:** There is no efficient Graph Isomorphism test that can be applied to **any class** of graphs!
- **Solution**
  - Visual similarity
  - Graph statistics similarity

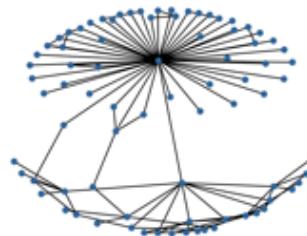
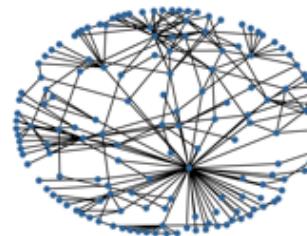
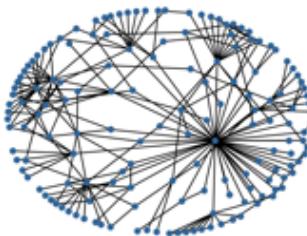
# Visual Similarity



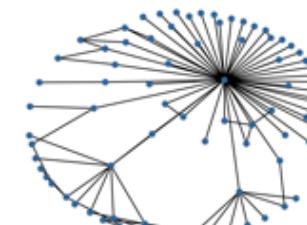
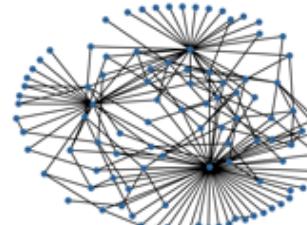
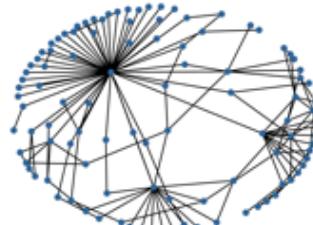
# Visual Similarity

Ego network

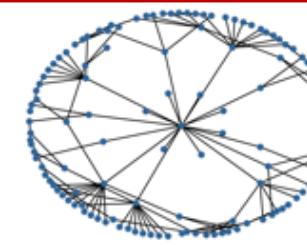
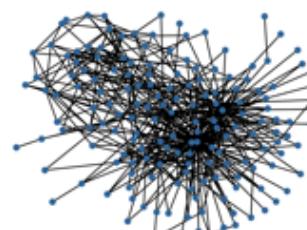
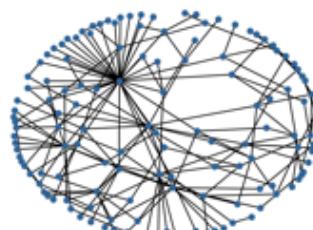
Training



GraphRNN



Baselines

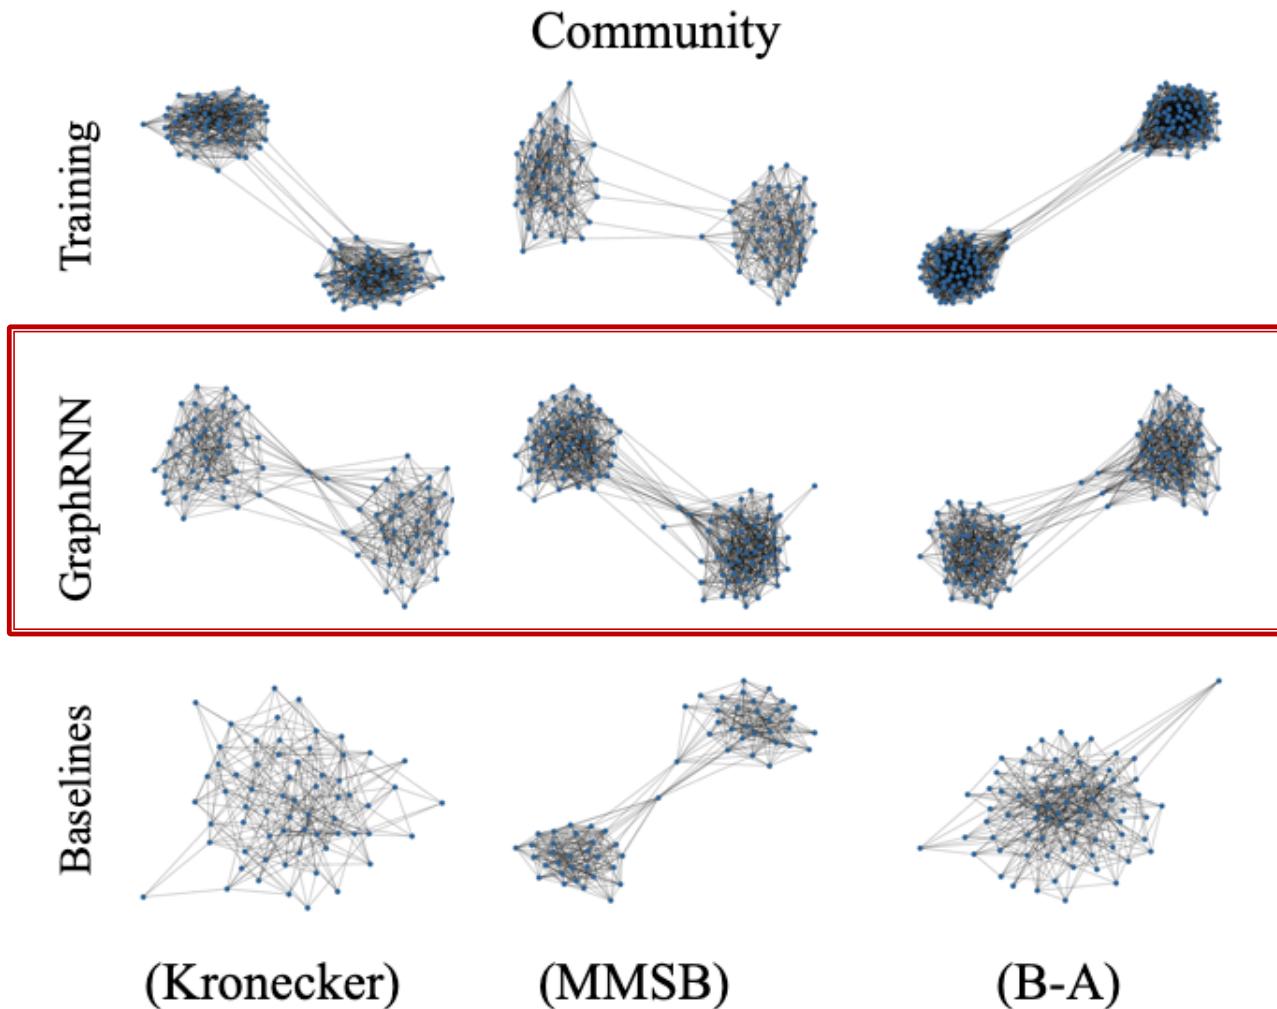


(Kronecker)

(MMSB)

(B-A)

# Visual Similarity



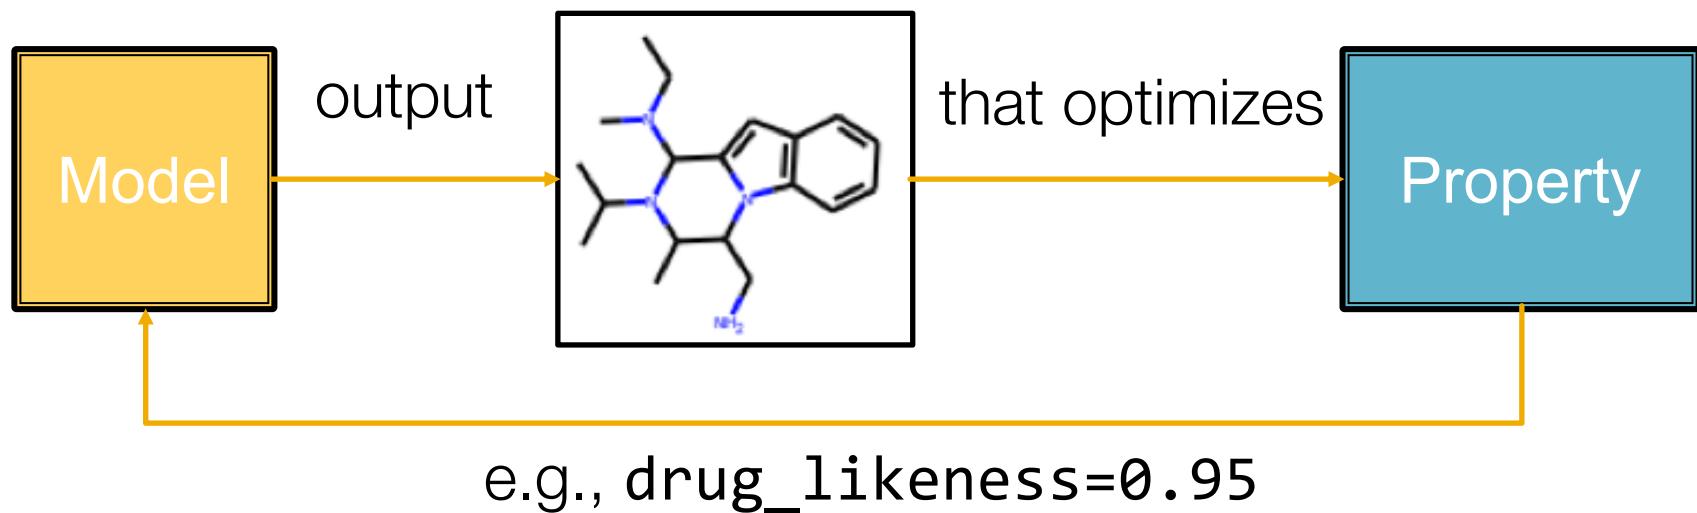
# Outline of Today's Lecture

- 1. Problem of Graph Generation**
- 2. ML Basics for Graph Generation**
- 3. GraphRNN**
- 4. Applications and Open Questions**



# Application: Drug Discovery

**Question:** Can we learn a model that can generate **valid** and **realistic** molecules with **high value of a given chemical property**?



[Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation](#). J. You, B. Liu, R. Ying, V. Pande, J. Leskovec. *Neural Information Processing Systems (NeurIPS)*, 2018.

# Goal-Directed Graph Generation

## Generating graphs that:

- Optimize a given objective (High scores)
  - e.g., drug-likeness (black box)
- Obey underlying rules (Valid)
  - e.g., chemical valency rules
- Are learned from examples (Realistic)
  - e.g., Imitating a molecule graph dataset

[Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation](#). J. You, B. Liu, R. Ying, V. Pande, J. Leskovec. *Neural Information Processing Systems (NeurIPS)*, 2018.

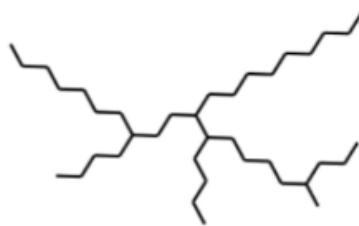
# Model: GCPN

**Graph Convolutional Policy Network** combines graph representation + RL:

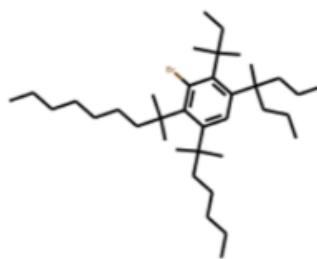
- **Graph Neural Network** captures complex structural information, and enables validity check in each state transition (**Valid**)
- **Reinforcement learning** optimizes intermediate/final rewards (**High scores**)
- **Adversarial training** imitates examples in given datasets (**Realistic**)

# Qualitative Results

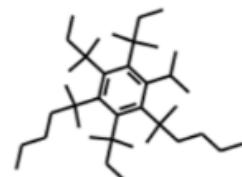
Visualization of GCPN graphs:  
Generate graphs with high property scores



7.98

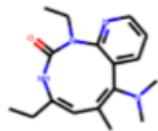


7.48

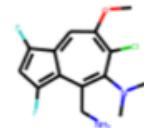


7.12

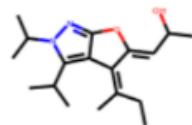
(a) Penalized logP optimization



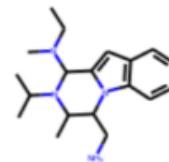
0.948



0.945



0.944



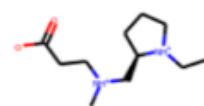
0.941

(b) QED optimization

# Qualitative Results

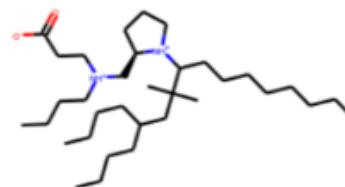
Visualization of GCPN graphs:  
Edit given graph for higher property scores

Starting structure

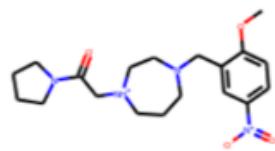


-8.32

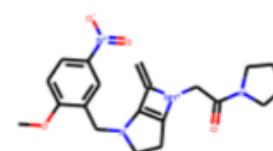
Finished structure



-0.71



-5.55



-1.78

(c) Constrained optimization of penalized logP

# Open Problems

- Generating graphs in other domains
  - 3D shapes, point clouds, scene graphs, etc.
- Scale up to large graphs
  - Hierarchical action space, allowing high-level action like adding a structure at a time
- Other applications: Anomaly detection
  - Use generative models to estimated prob. of real graphs vs. fake graphs

# Outline of Today's Lecture

- 1. Problem of Graph Generation** 
- 2. ML Basics for Graph Generation** 
- 3. GraphRNN** 
- 4. Applications and Open Questions** 